

Realistic / Advanced Image Synthesis

Tutorial Week #2

Fall 2019

Topics for today

- I. Camera view
- II. Rendering loops
- III. Getting started with the assignments



Teaching assistants



Joey

Mon & Tue*

{joey.litalien, peter.quinn}@mail.mcgill.ca

Peter

Wed & Fri*

Tutorials

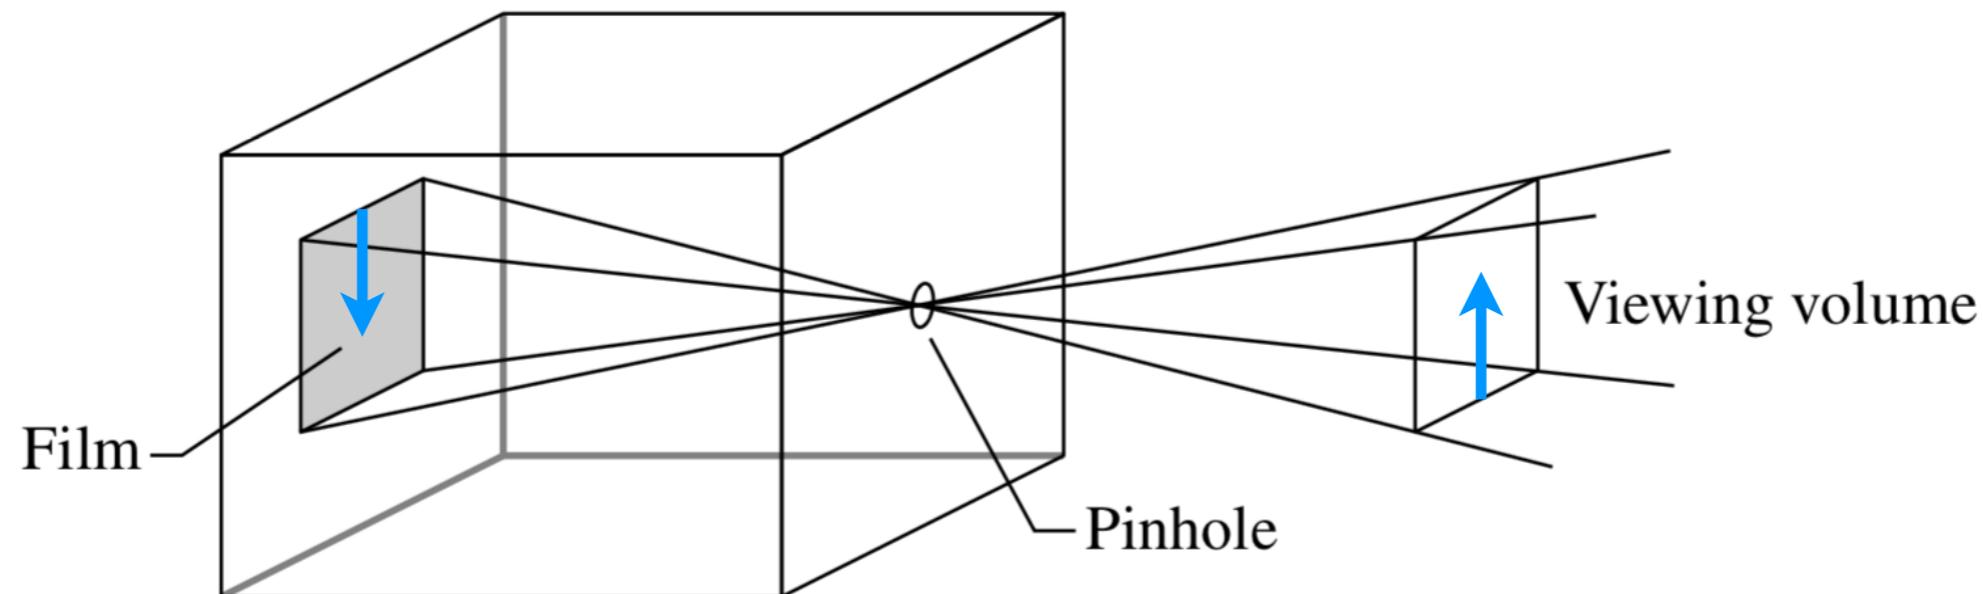
- ▶ Complement the lectures
- ▶ Answer your questions
- ▶ **Help you with the assignment code**
- ▶ But... don't expect us to magically solve your coding bugs!

Camera view



Pinhole camera

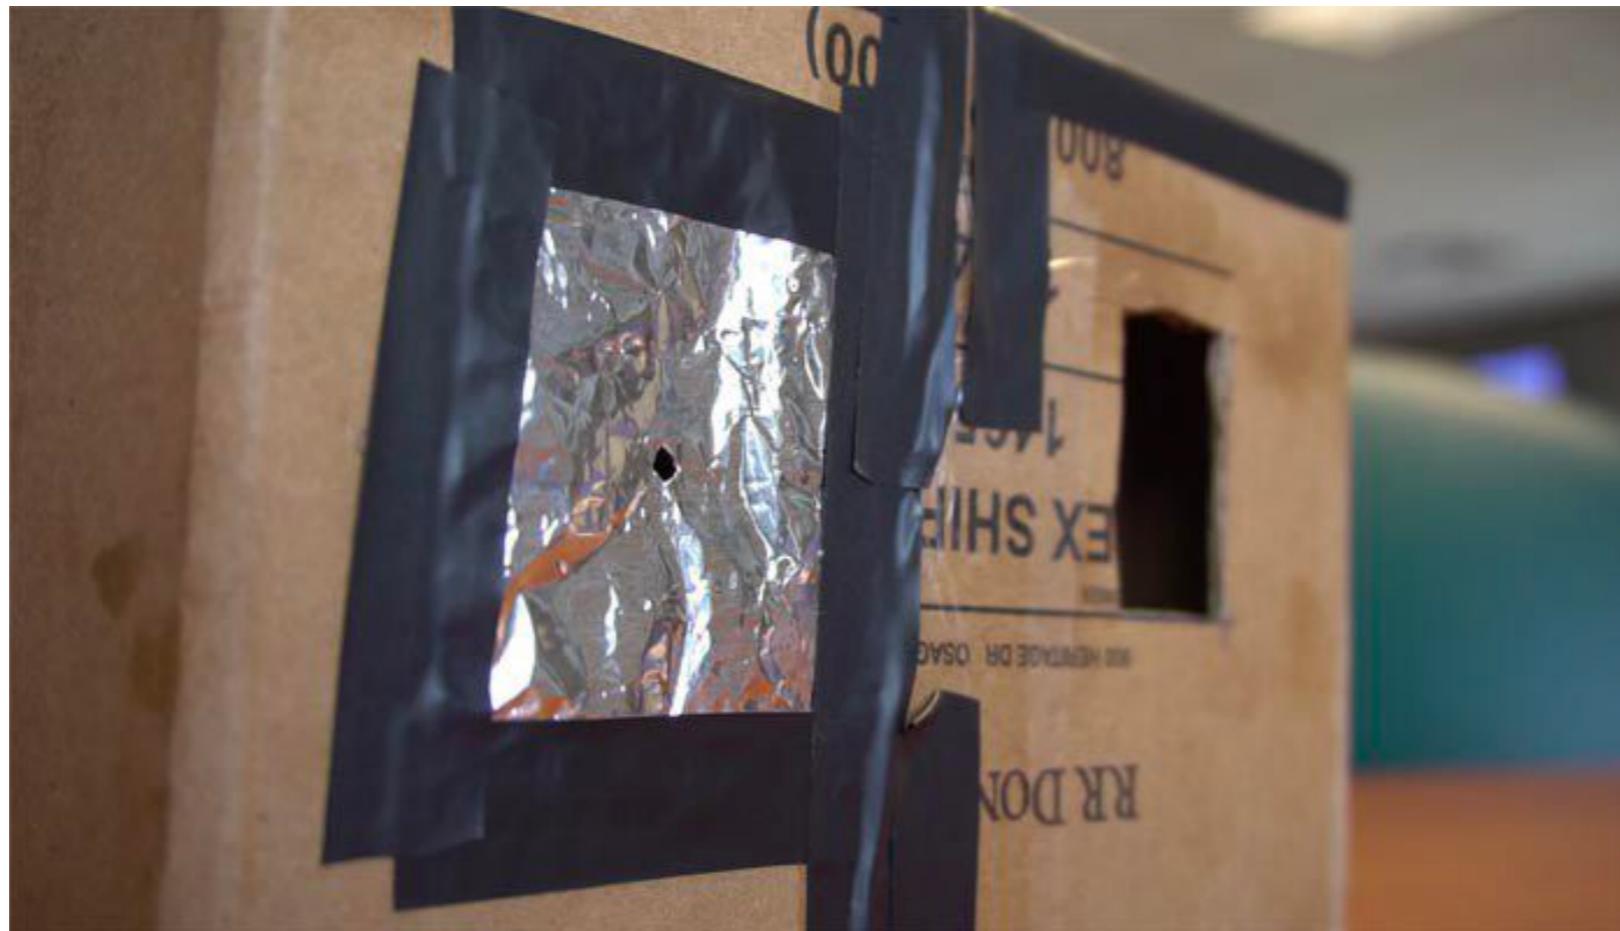
- ▶ Infinitely small aperture, hence the name
- ▶ Defined with position, up vector and look-at vector



Camera view

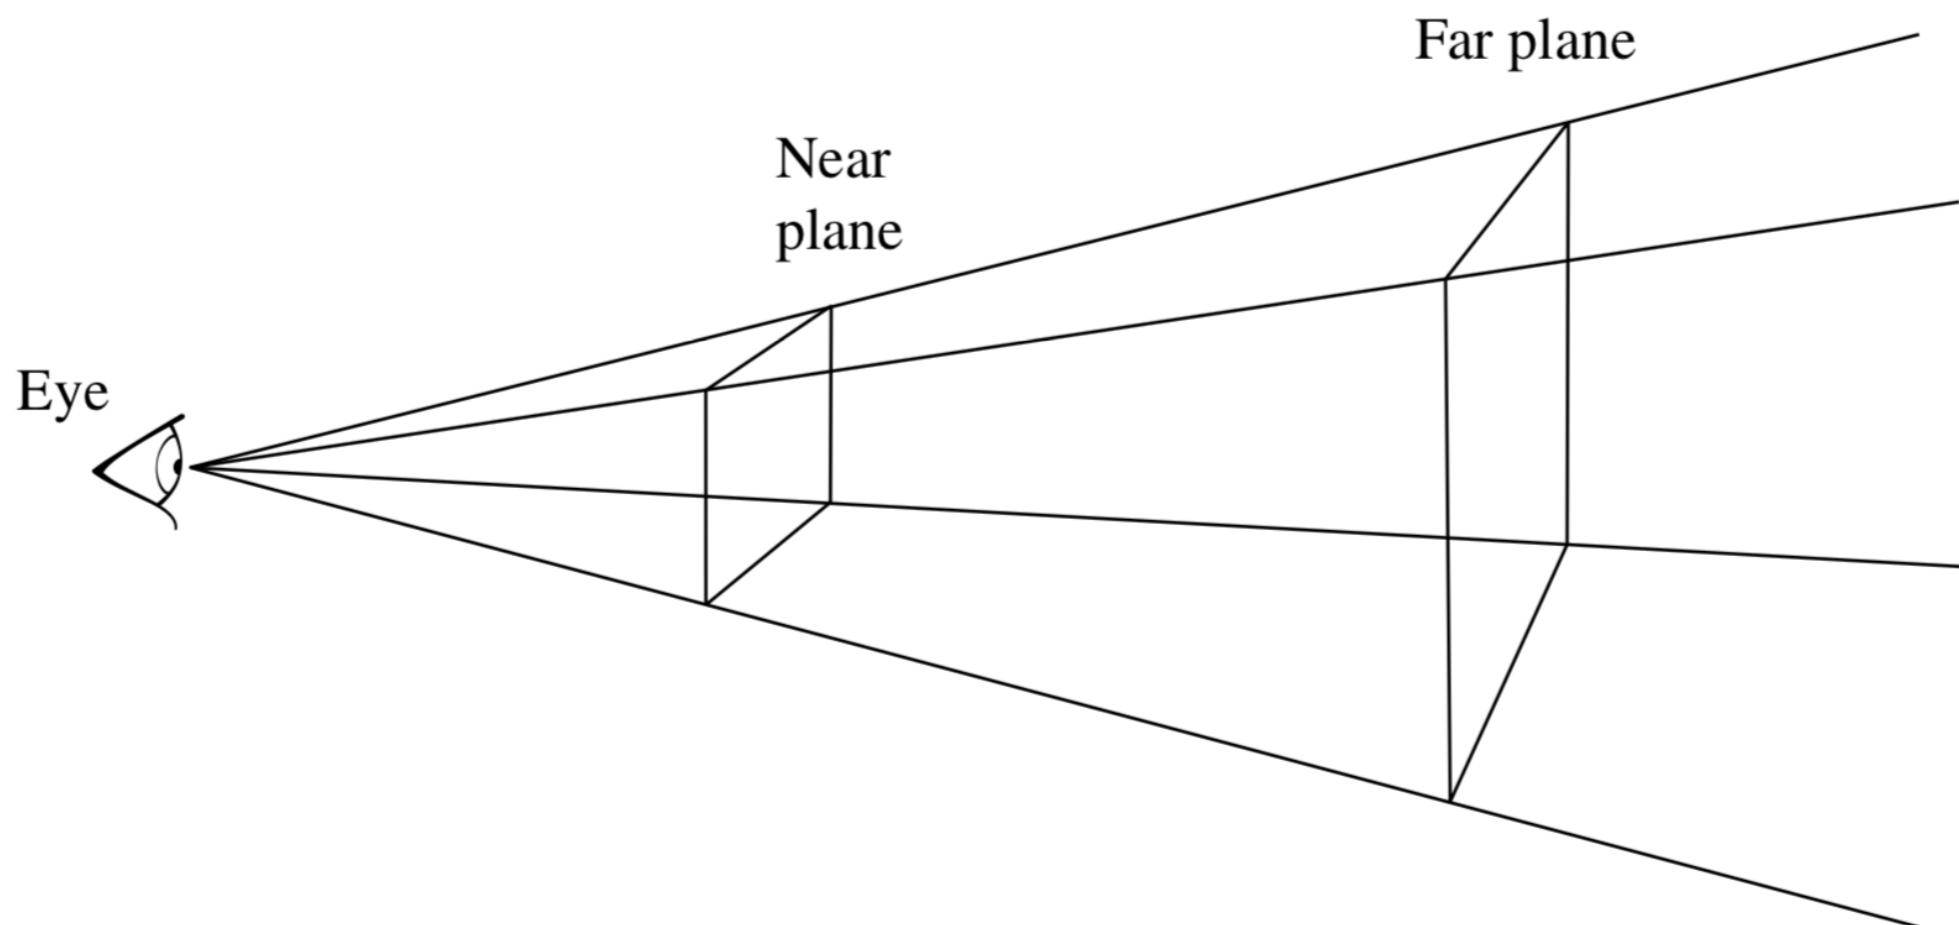
Pinhole camera IRL

- ▶ Hole in a cardboard box to see solar eclipses



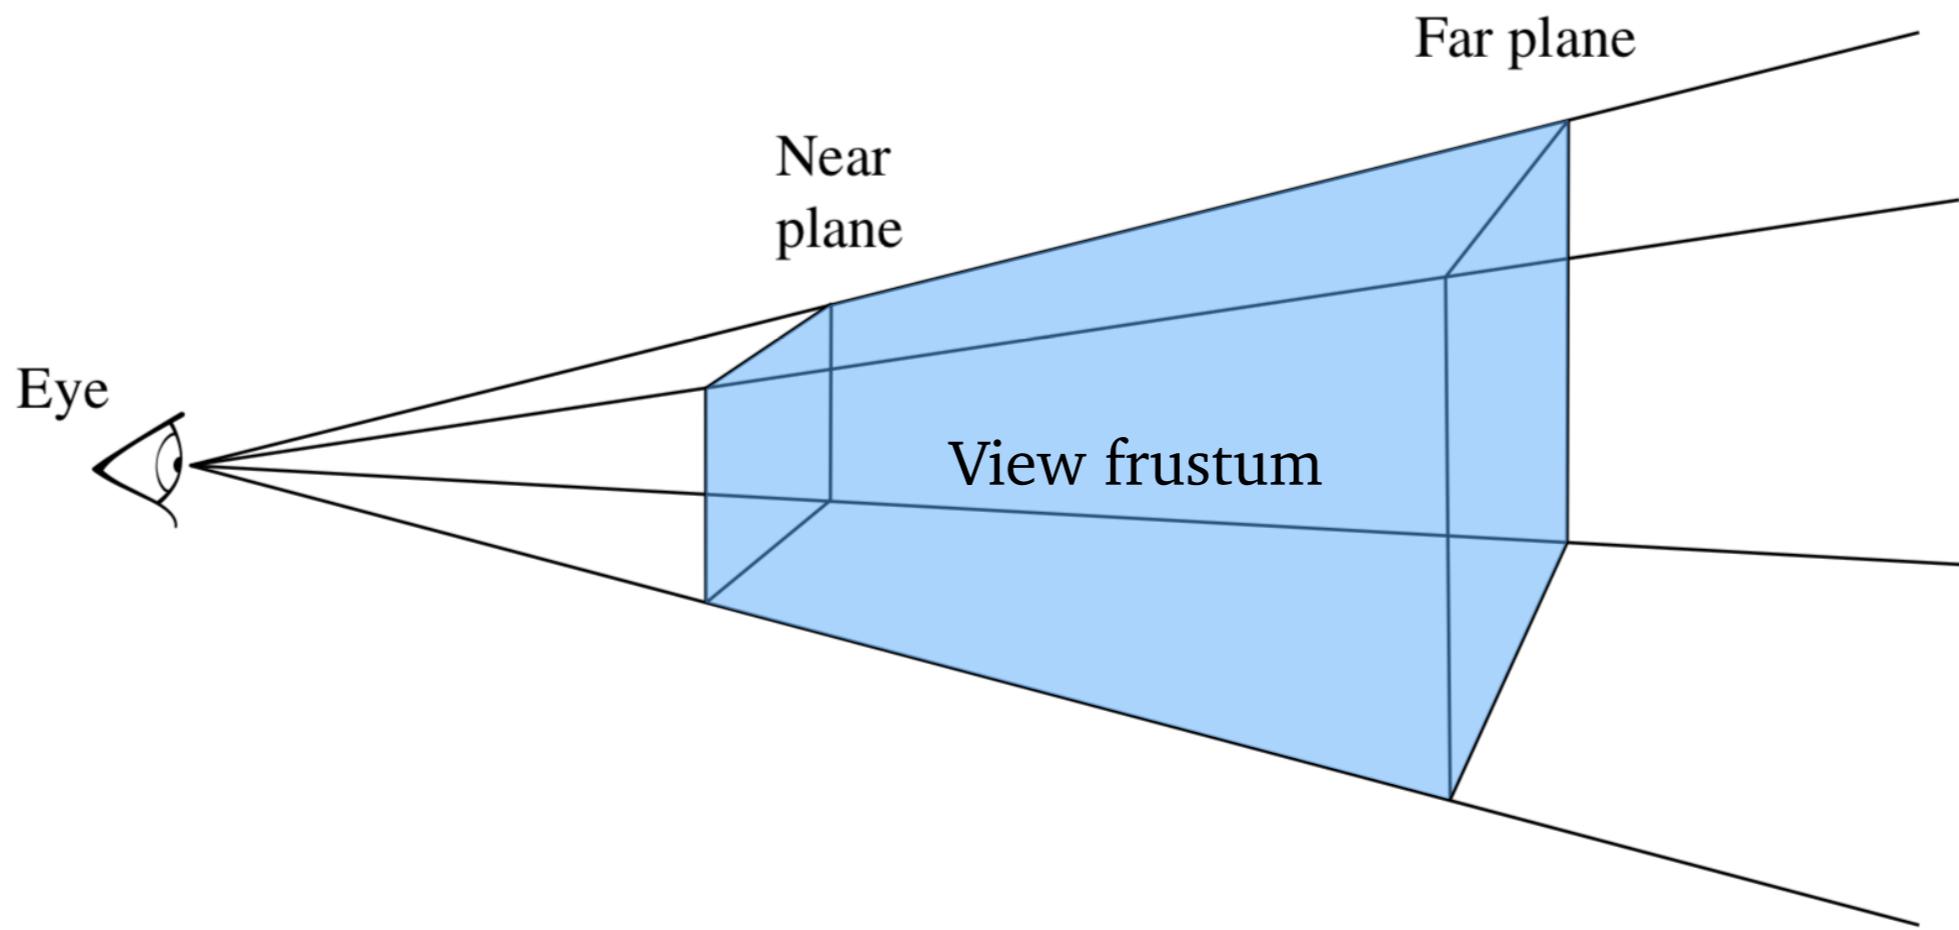
Pinhole camera

- ▶ Simulate pinhole by placing film *in front of* hole
- ▶ *Hole* is what we call *eye* in rendering



Pinhole camera

- ▶ Simulate pinhole by placing film *in front of* hole
- ▶ *Hole* is what we call *eye* in rendering



Coordinate spaces

- ▶ **Object space**

Coordinate system in which geometric primitives are defined

- ▶ **World space**

Standard frame that all other spaces are defined in terms of.
Each primitive has an object-to-world transformation

- ▶ **Camera space**

New coordinate system with origin at camera eye

Negative z -axis = Viewing direction

y -axis = Up direction

Camera matrices

- ▶ **Model matrix** (Identity in this course)

Acts on the vertices of the object in local coordinates.

Model = Translation · Rotation · Scale (of your actual mesh)

- ▶ **View matrix**

World-to-camera transform that puts the camera at the center of the space

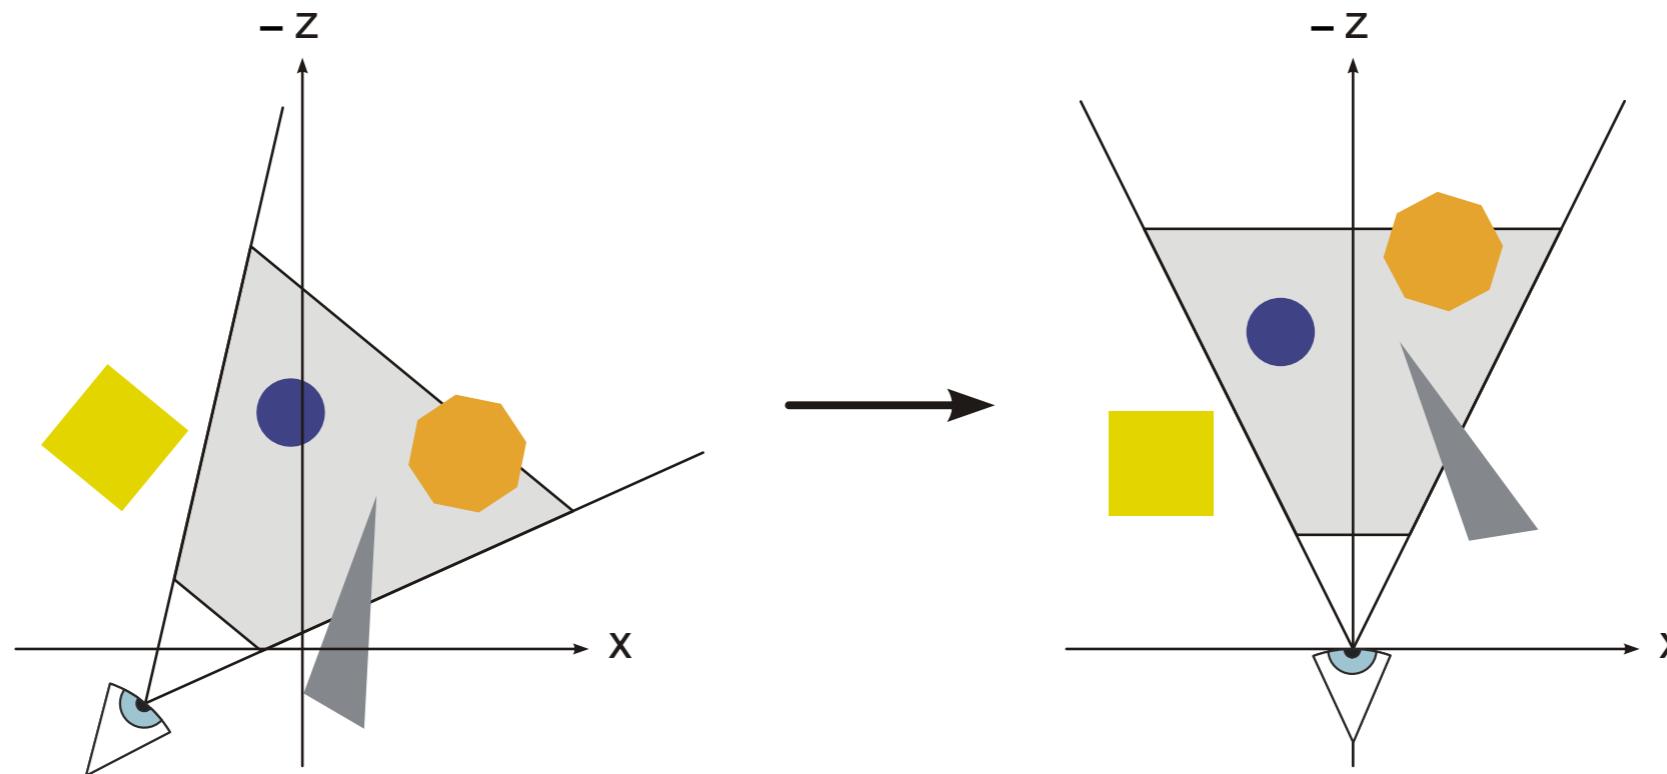
- ▶ **Projection matrix**

Perspective transform from camera space to clip coordinates.

Simulates how we perceive things IRL

View matrix

- ▶ Moves the world so that camera is at origin
- ▶ Expressed as a 4×4 transformation matrix

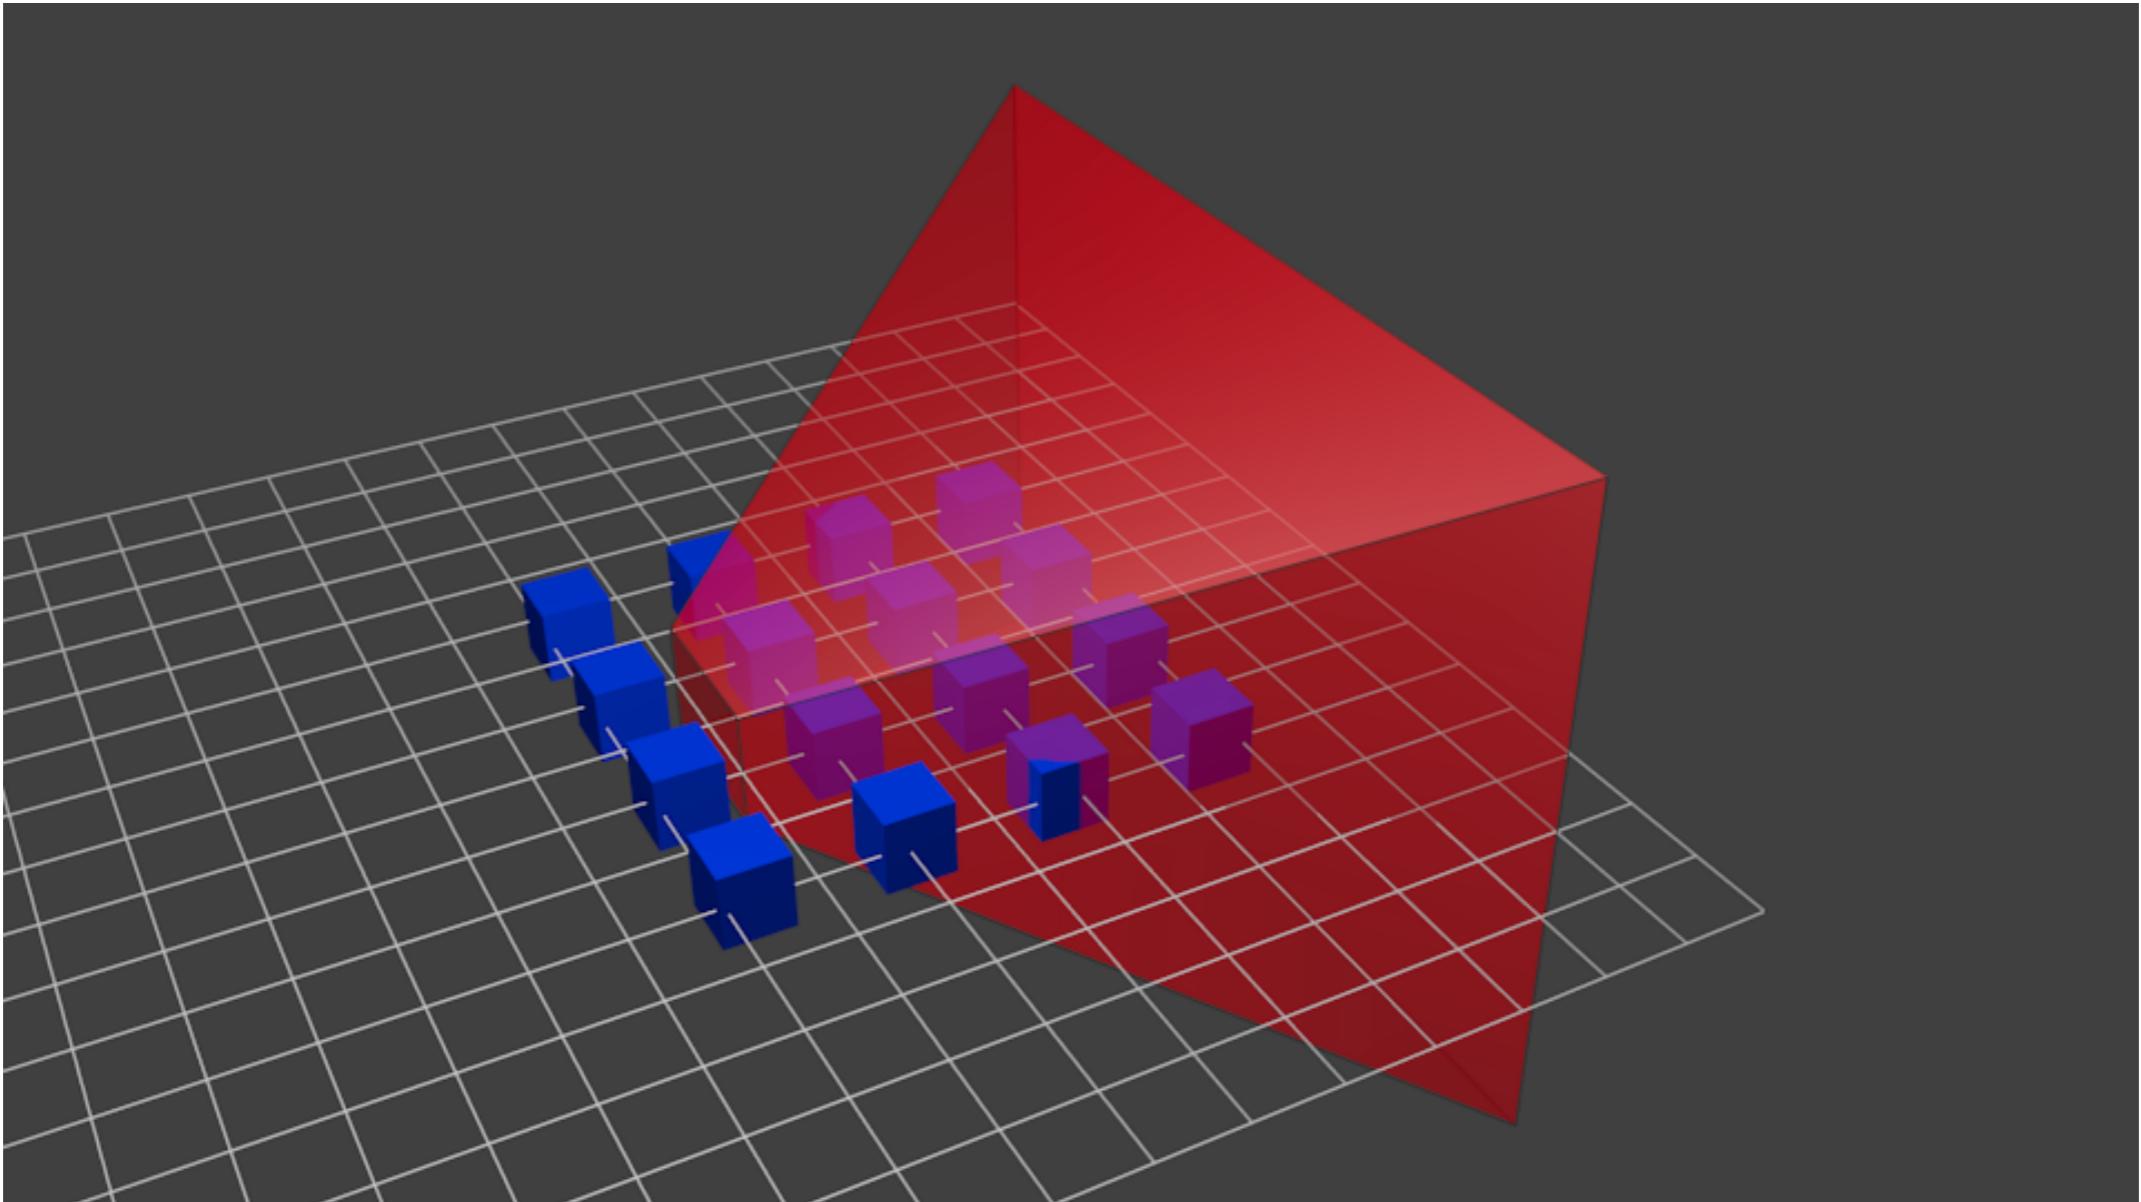


Projection matrix

- ▶ Transforms view frustum (truncated pyramid) into unit cube
- ▶ Expressed as a 4×4 projection matrix
- ▶ After transformation, models are said to be in *clip coordinates*

Camera view

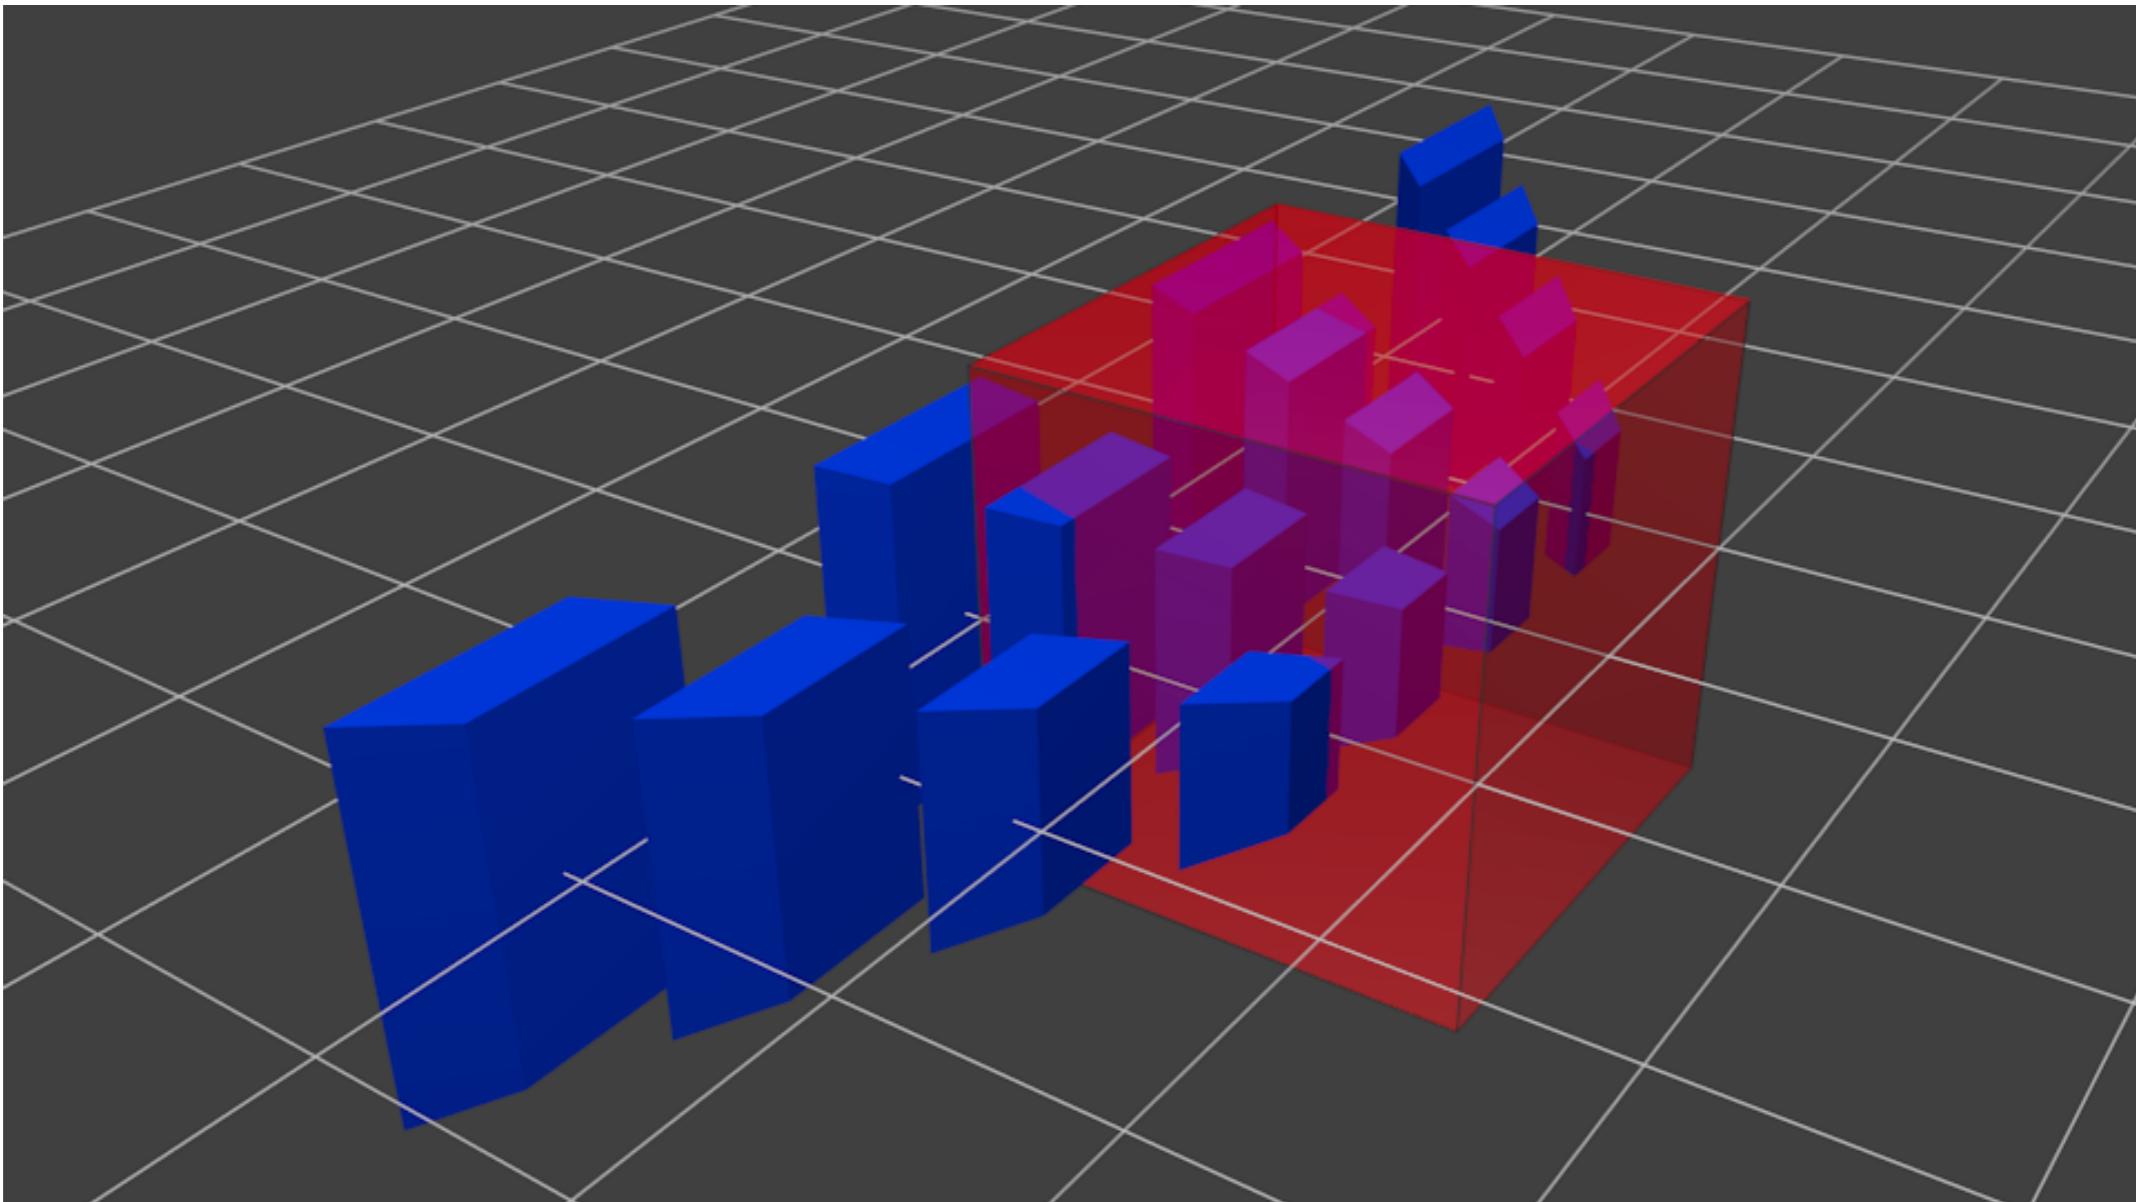
Projection matrix



Before (camera space)

Camera view

Projection matrix



After (clip coordinates)

Further readings

- ▶ **Offline**

PBRT3 Chapter 6.1 *Projective camera models*

- ▶ **Real-time**

RTR4 Chapter 4.7 *Projections*

Good tutorial (includes the previous images for the projection

matrix): [http://www.opengl-tutorial.org/
beginners-tutorials/tutorial-3-matrices](http://www.opengl-tutorial.org/beginners-tutorials/tutorial-3-matrices)

Offline rendering loop

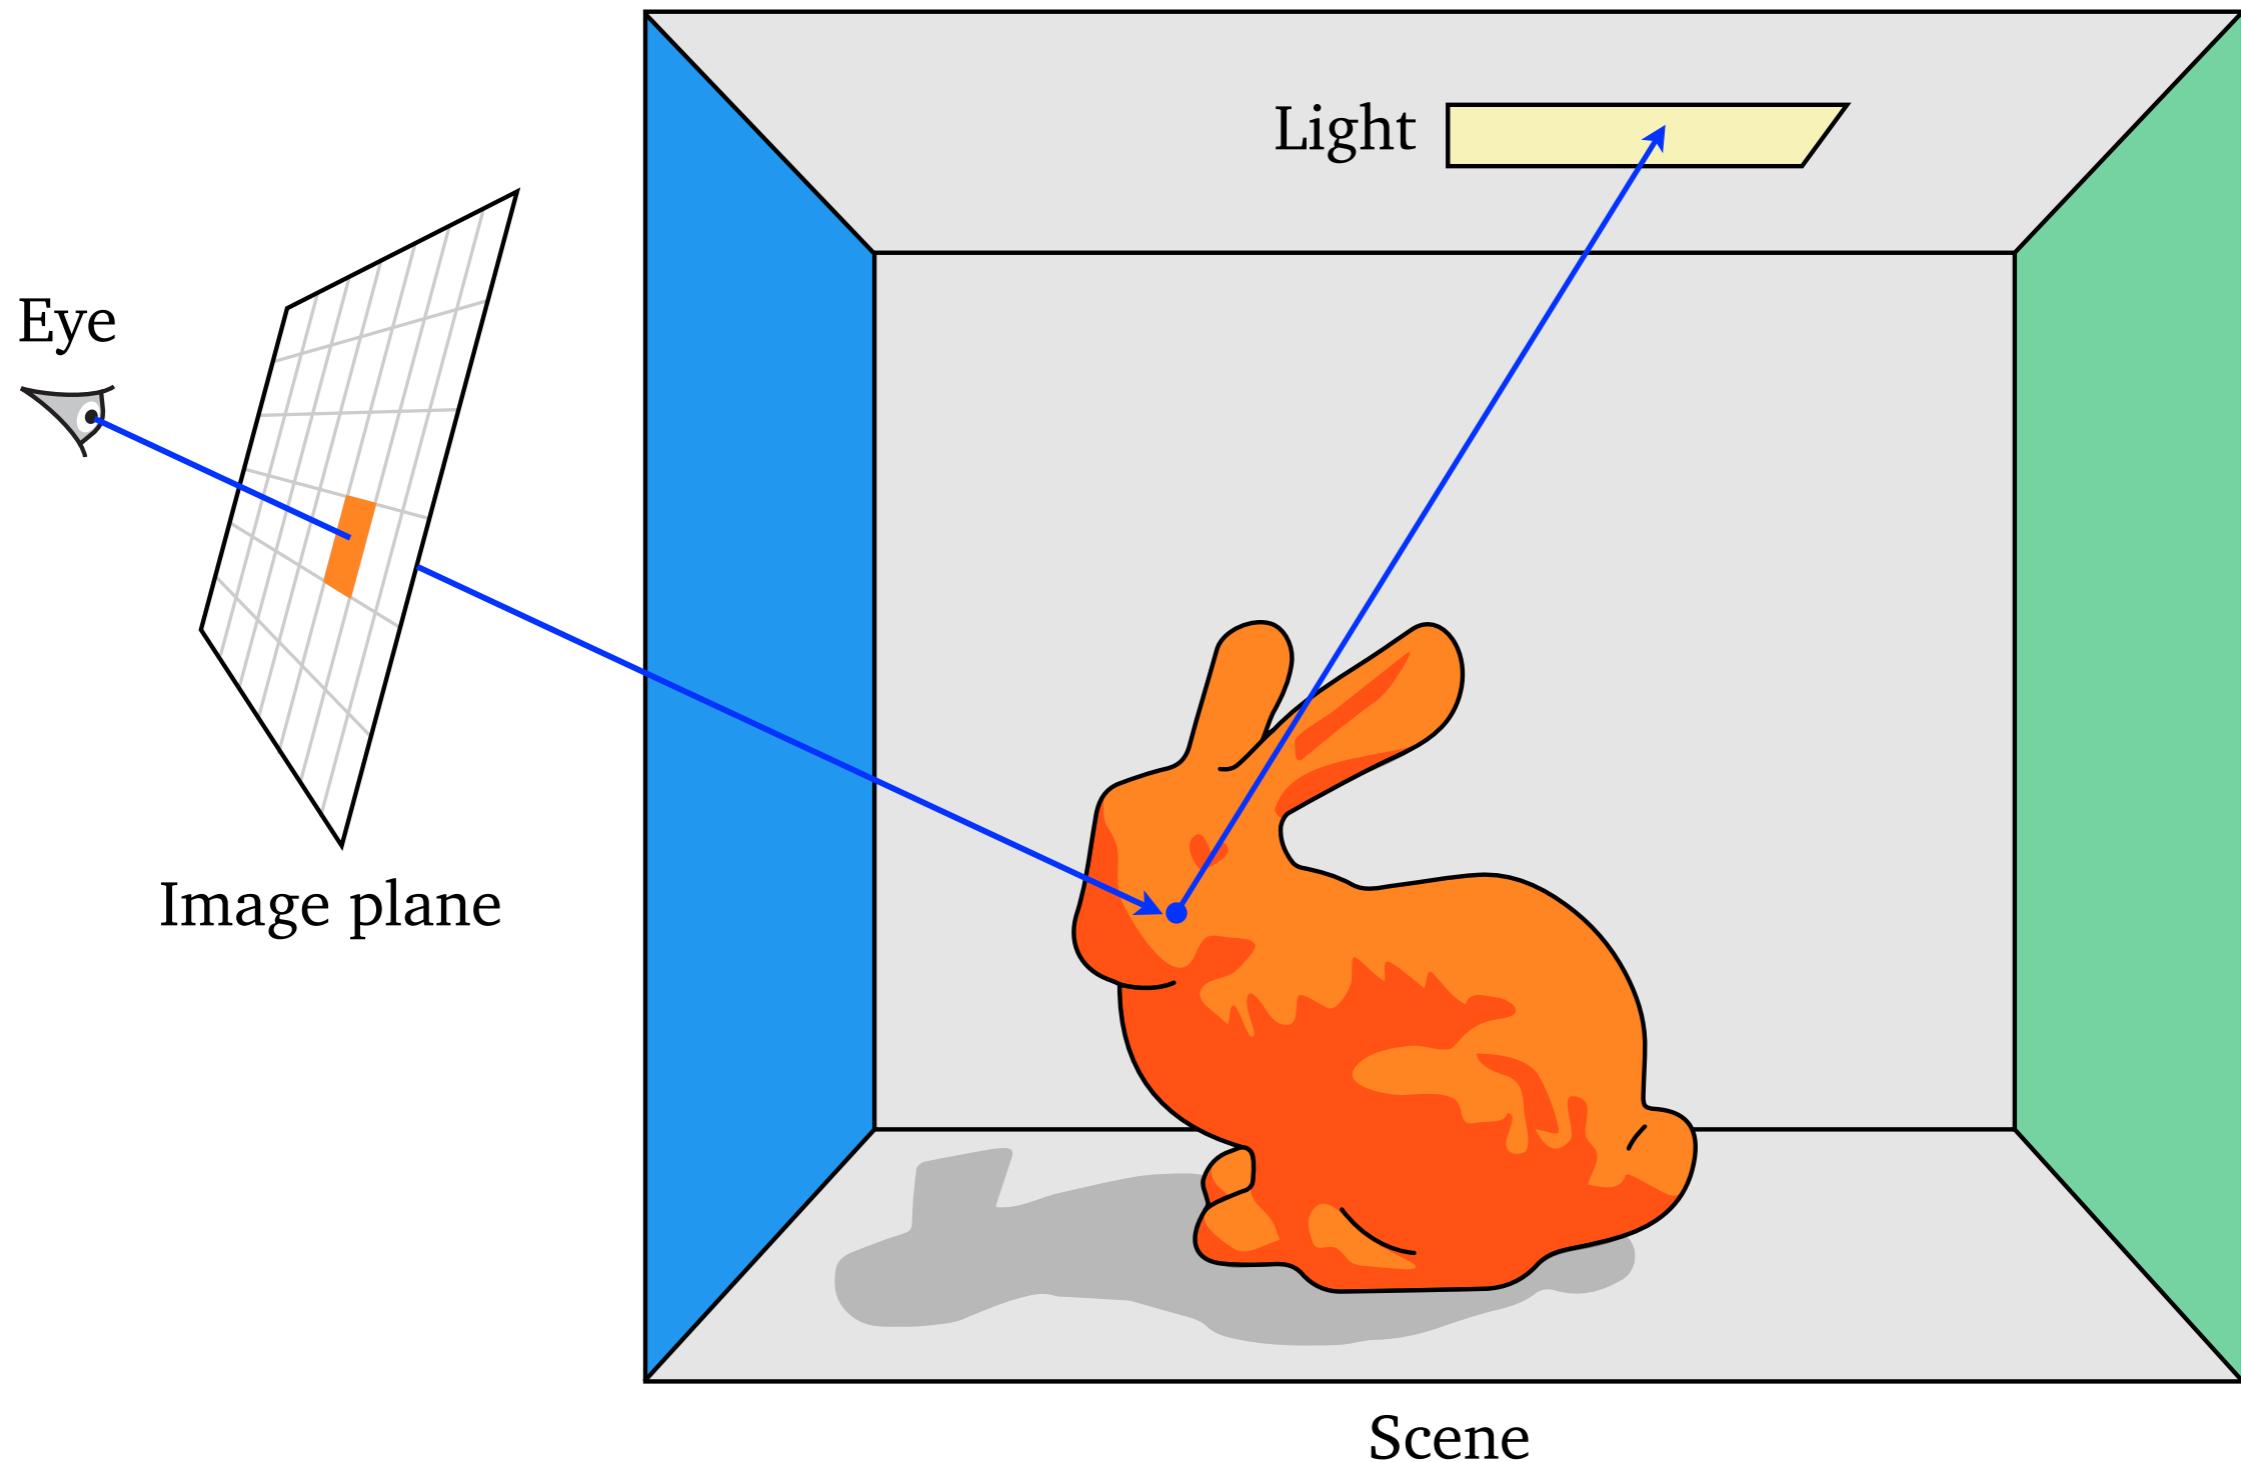


Ray tracing 101

1. Shoot rays from eye through each pixel of the image plane
2. Hit a surface, retrieve its reflectance property (normal, colour, etc.)
3. Evaluate light contribution (and possibly recurse)
4. ???
5. Profit

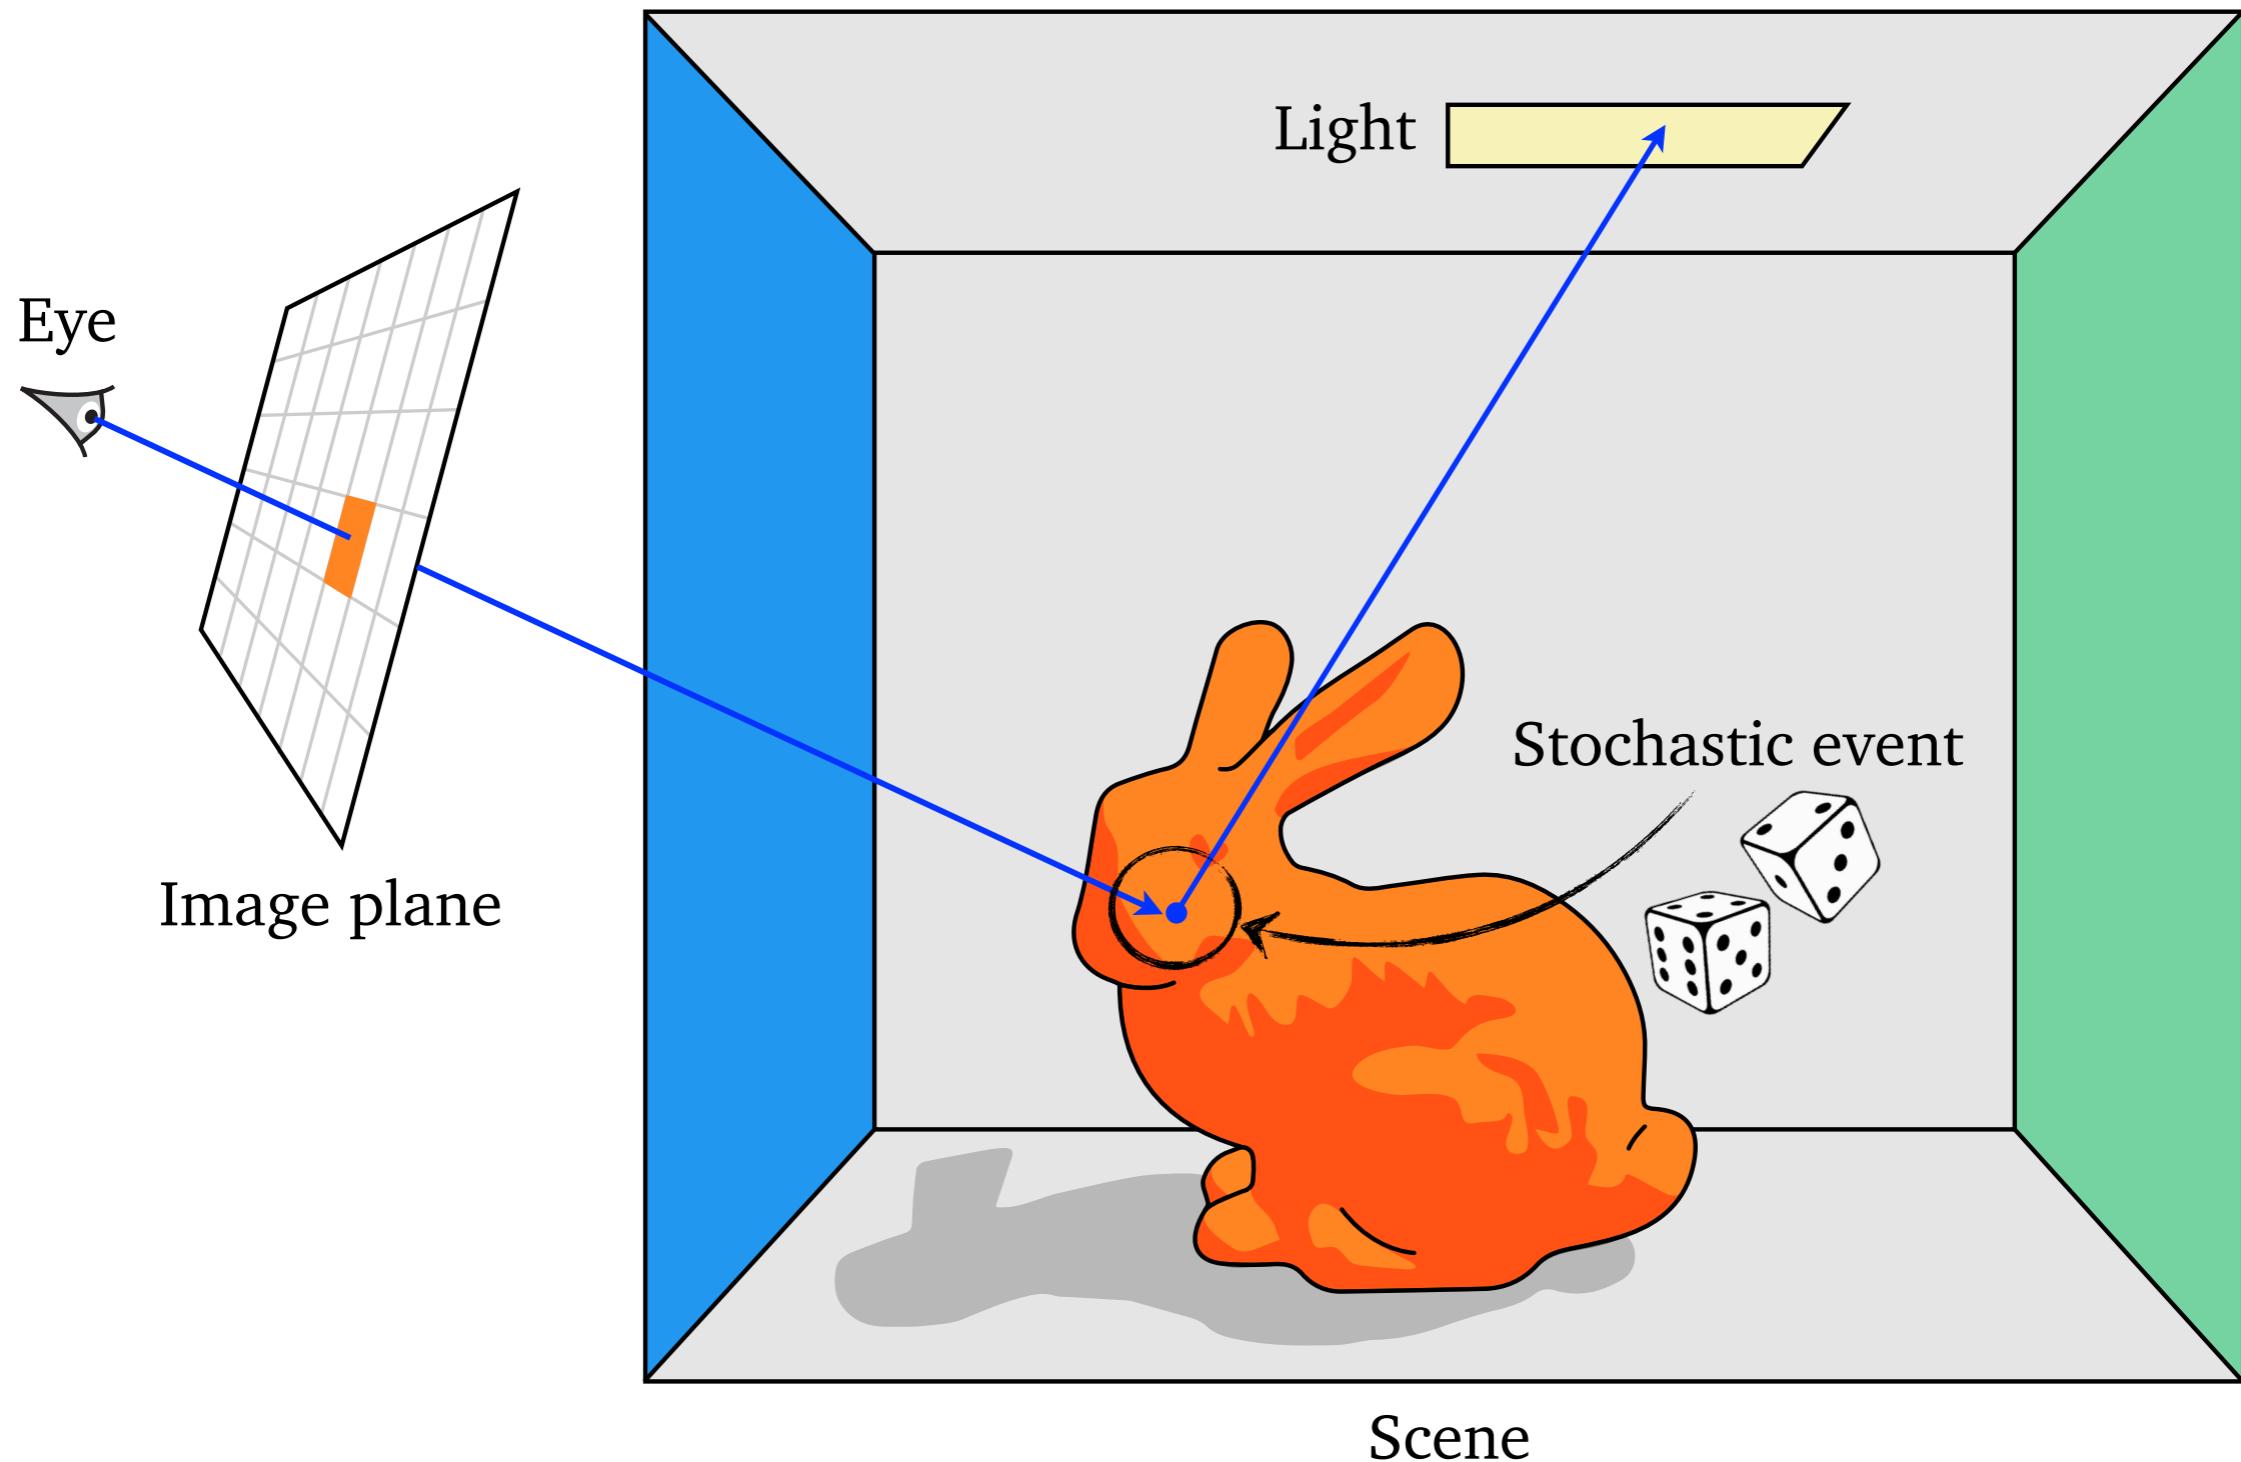
Offline rendering loop

Ray tracing 101



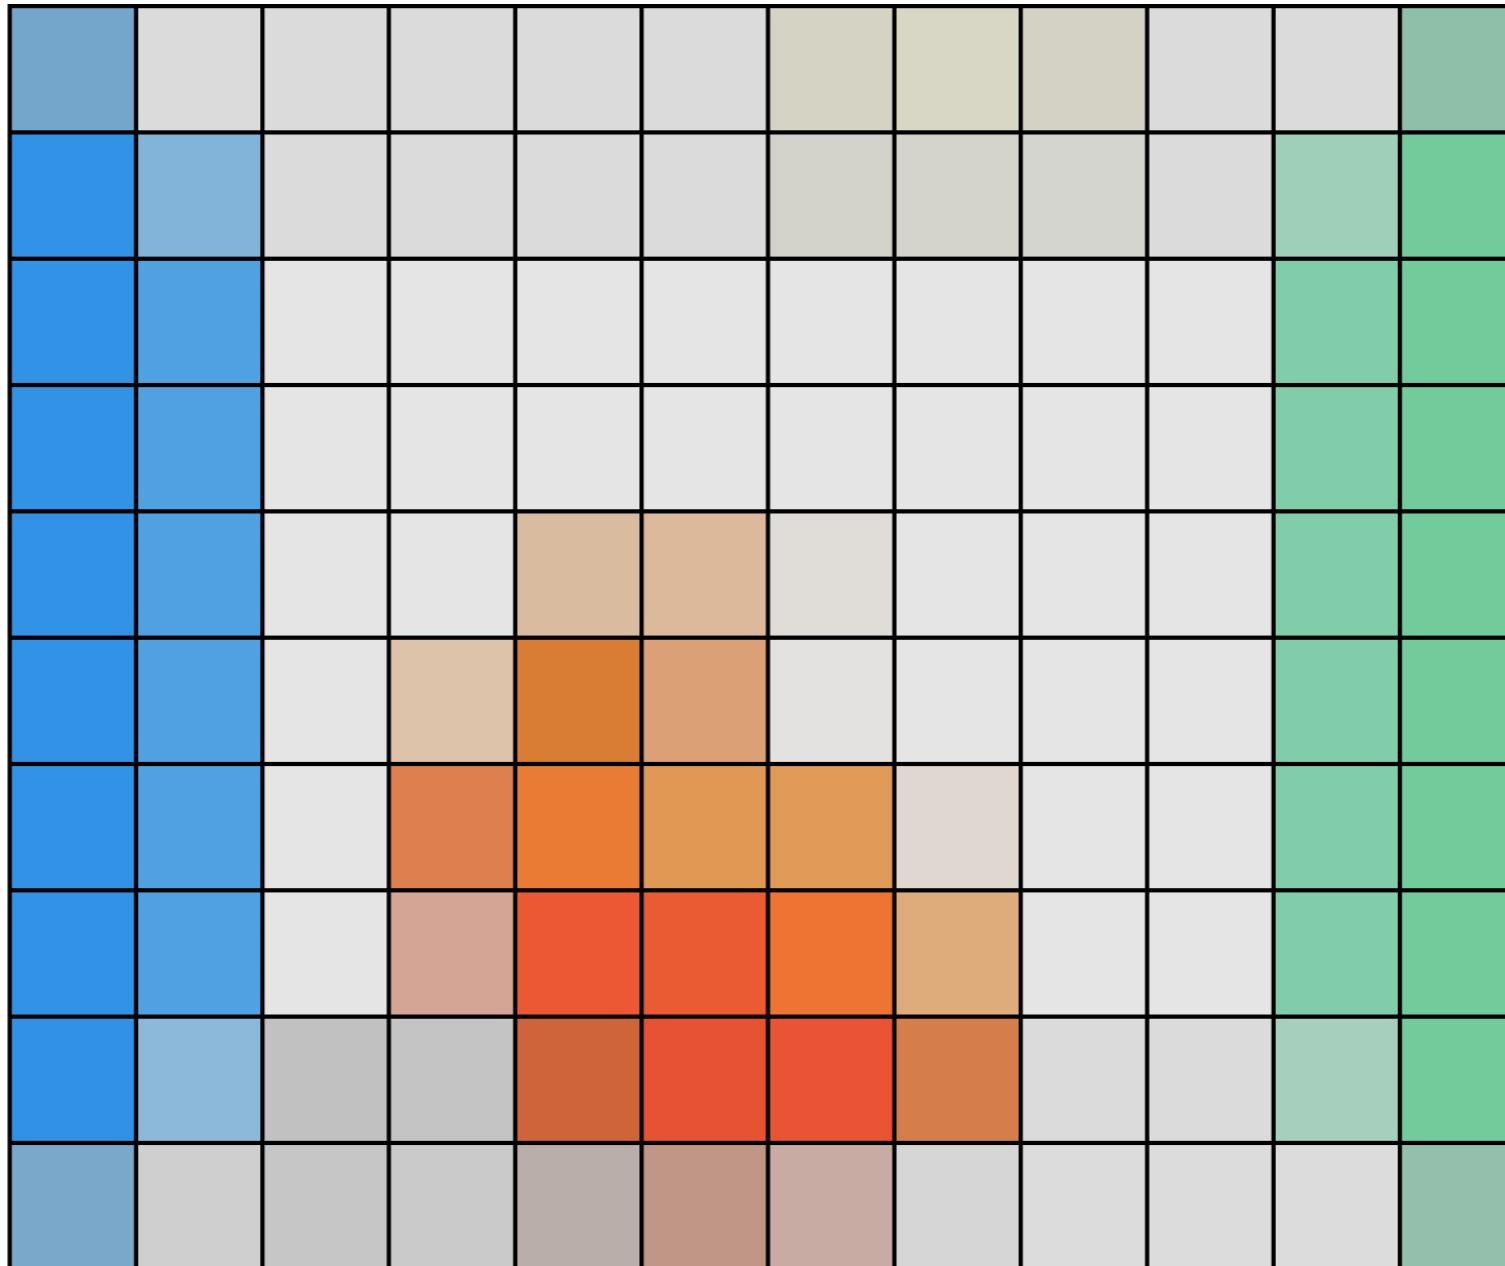
Offline rendering loop

Ray tracing 101



Offline rendering loop

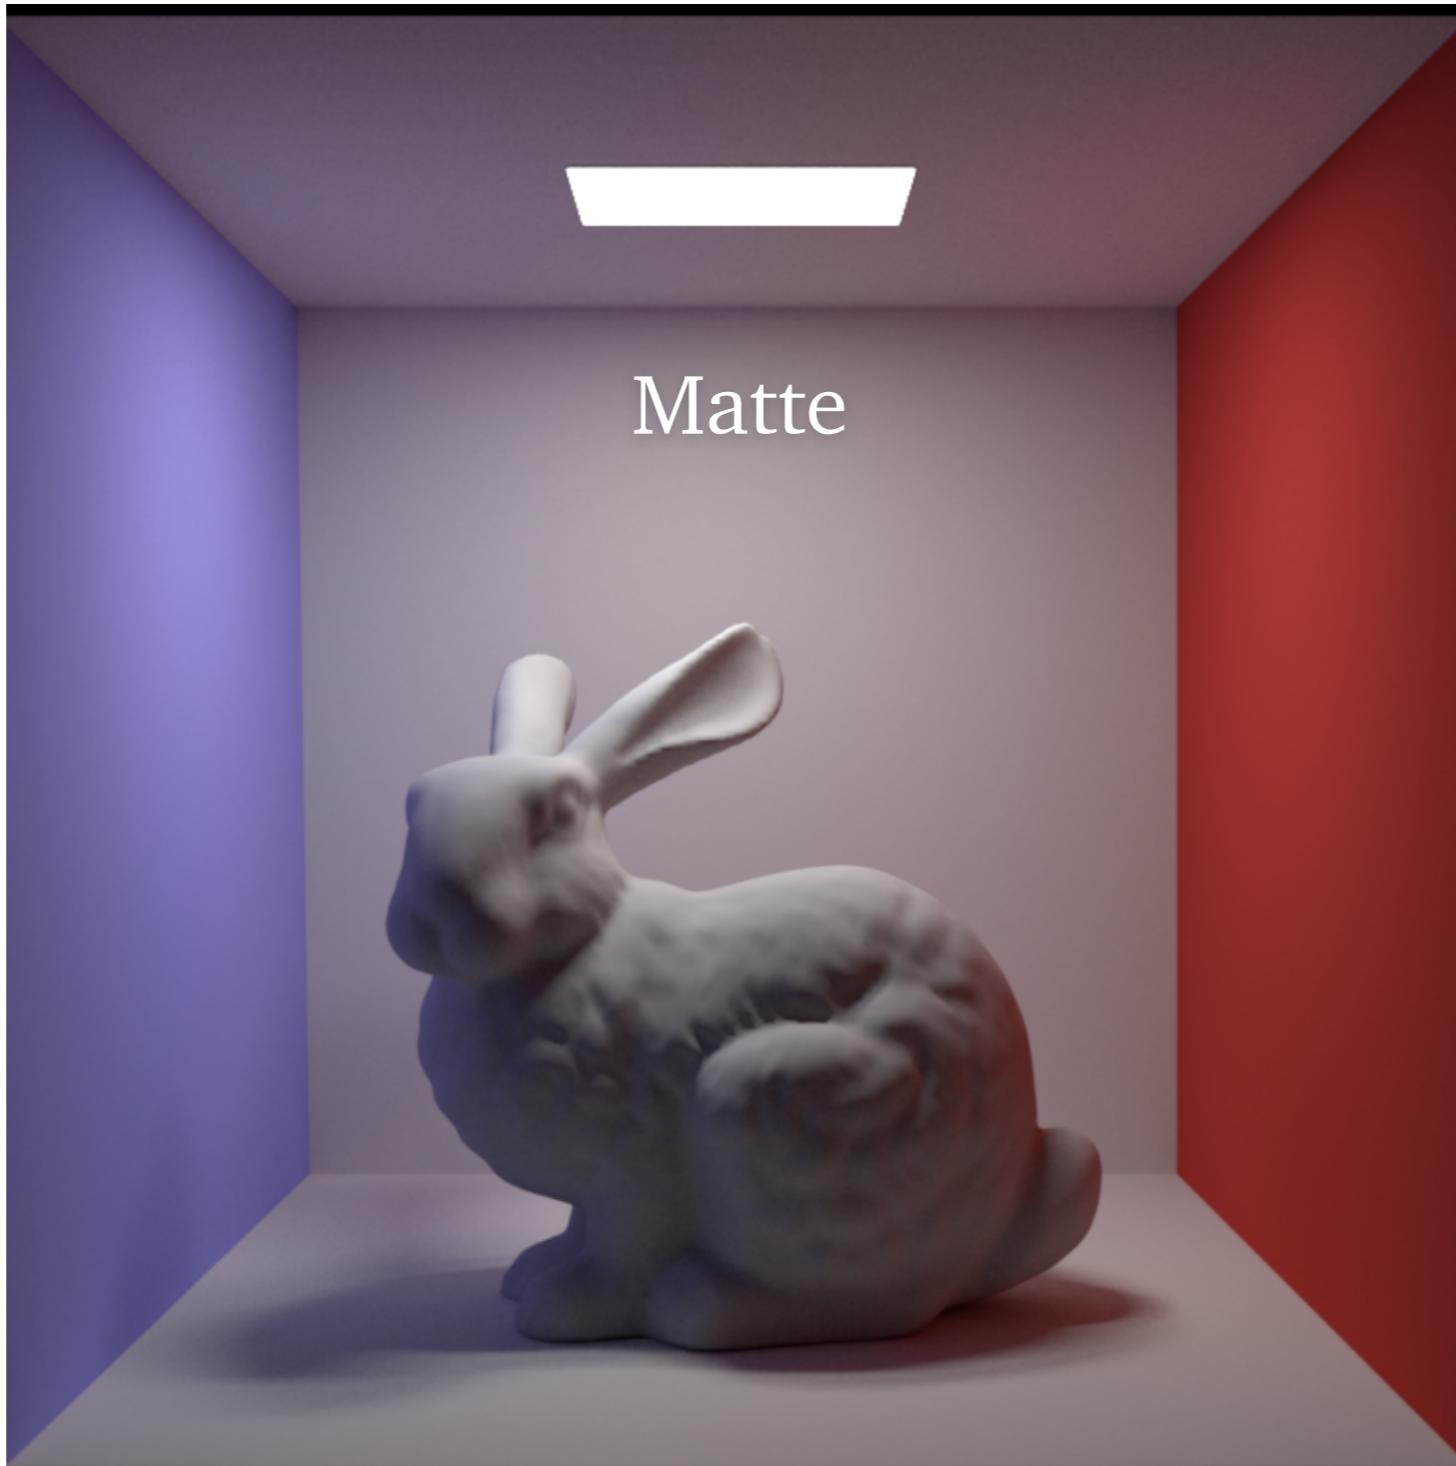
Ray tracing 101



Low-res image (after ray traced)

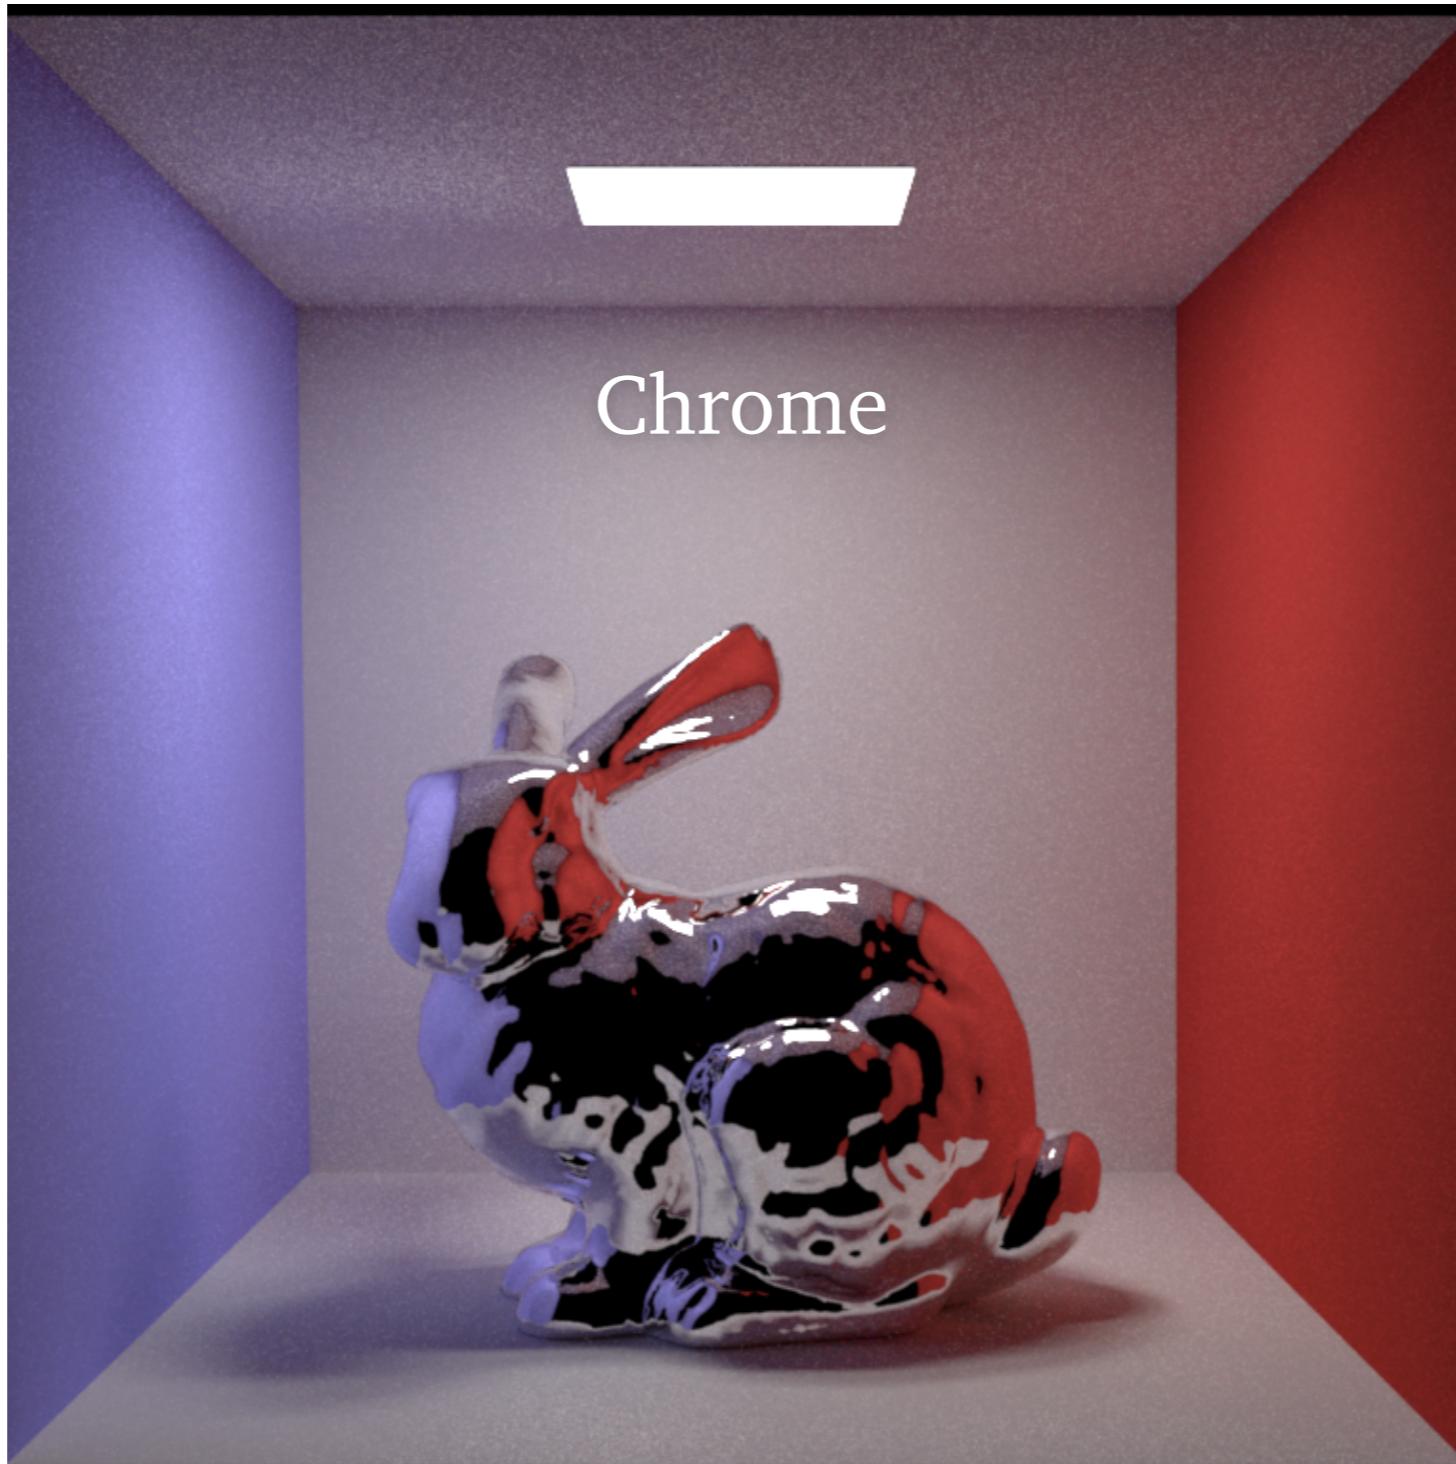
Offline rendering loop

Ray tracing 101



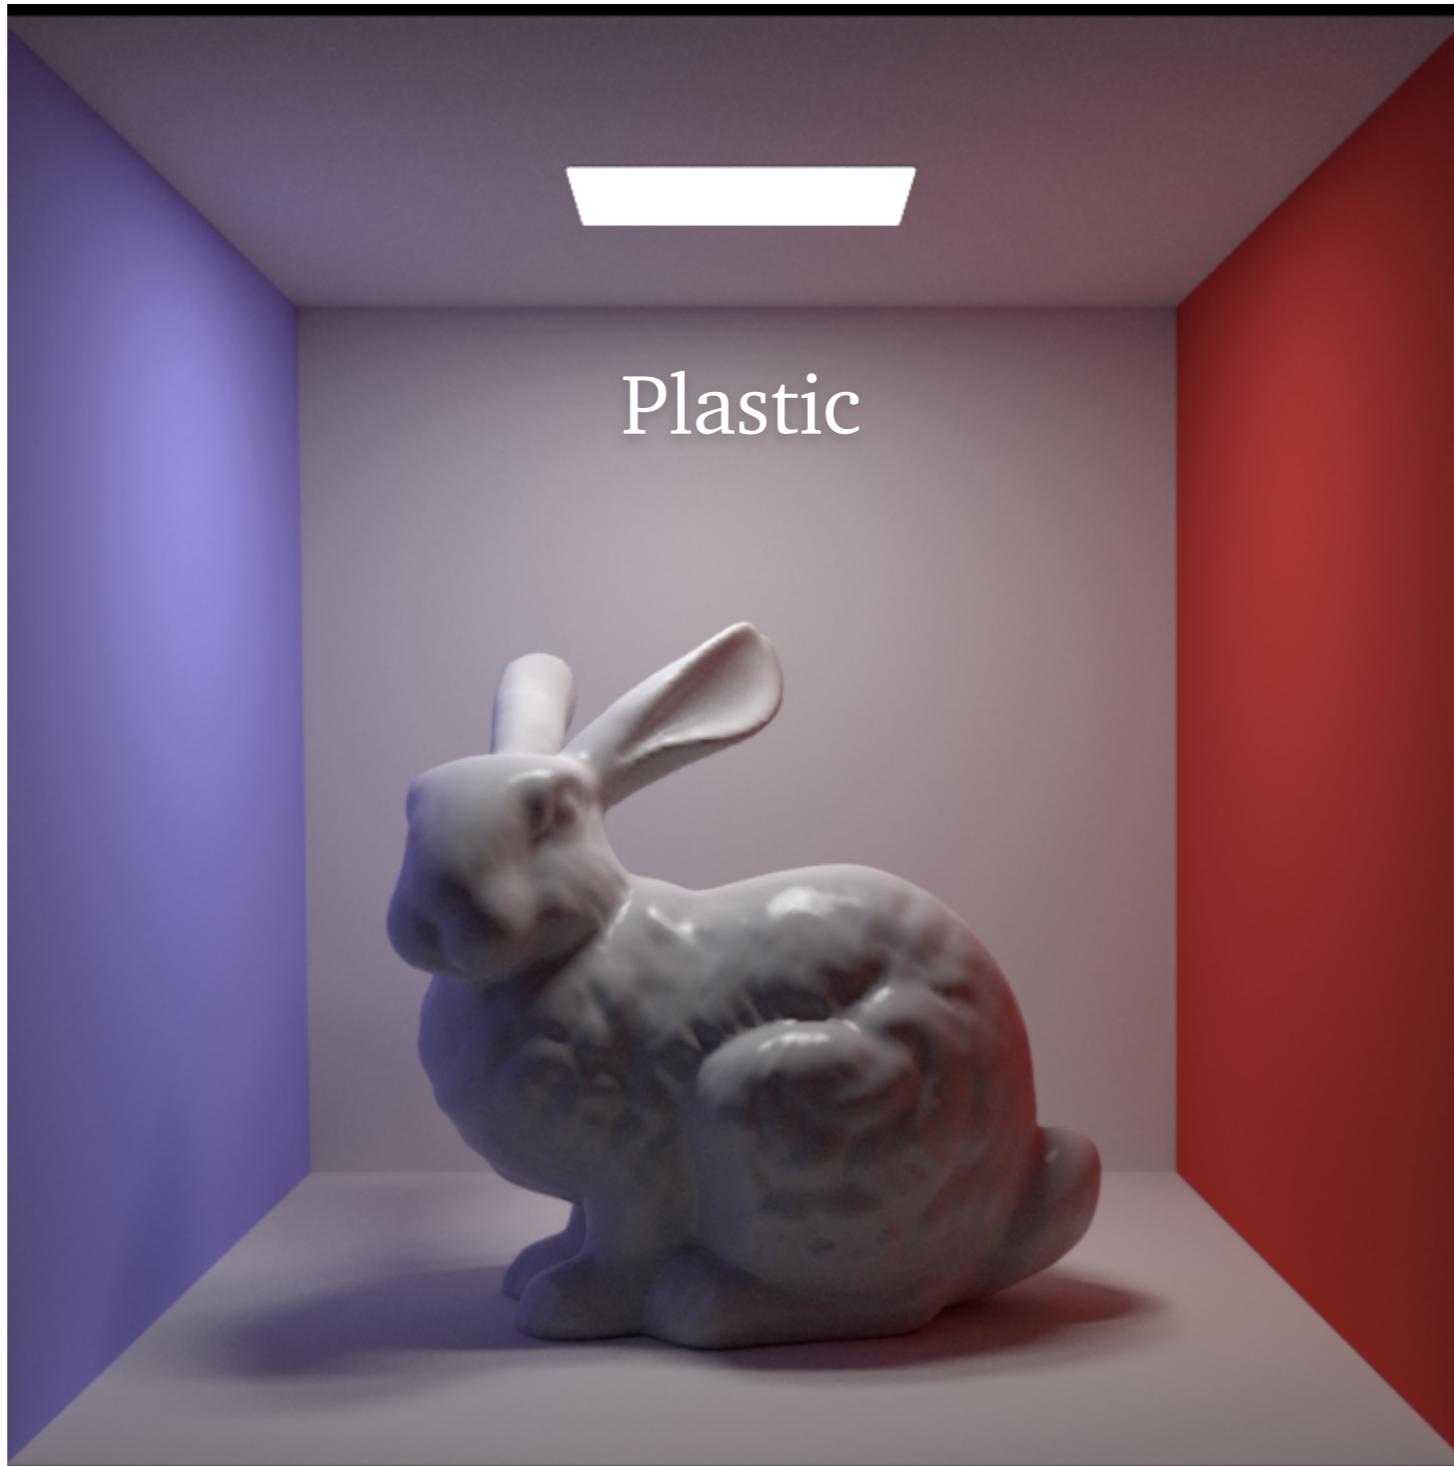
Offline rendering loop

Ray tracing 101



Offline rendering loop

Ray tracing 101



Offline rendering loop

Ray tracing 101



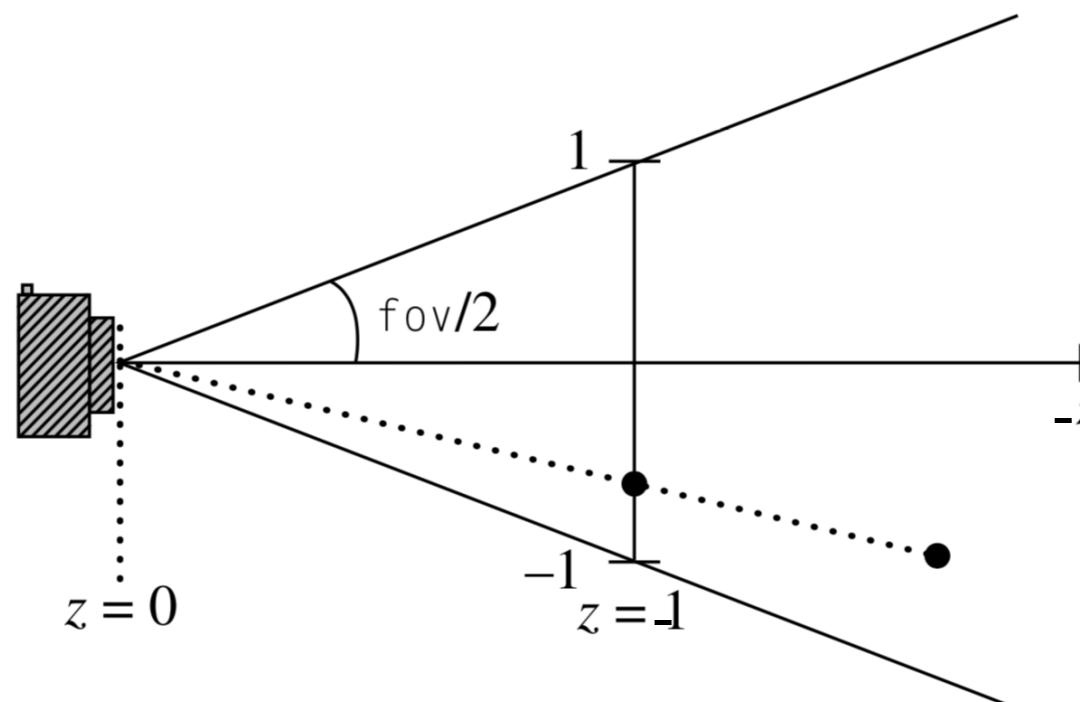
Field of view

- Pixel location should be in the interval $[-1, 1]$

- Account for FOV by scaling the pixel location:

```
scaling = tan(deg2rad(fov) / 2.f);
```

- Adjust film dimension by scaling x -axis by aspect ratio (w/h)



Field of view

- ▶ TinyRender uses *vertical* FOV
- ▶ Seems counterintuitive as we usually prefer landscape mode in photography
- ▶ Arbitrary choice to match camera settings of reference scenes, so just roll with it

View matrix

- ▶ Use OpenGL Mathematics:

```
glm::mat4 inverseView =  
    glm::inverse(glm::lookAt(eye, at, up));
```

- ▶ Gives you the matrix in *row major order*
- ▶ To map a ray from camera to world spaces:
 1. Augment to homogeneous coordinates (4D)
 2. Multiply by inverse view matrix *on the left*

Offline rendering loop

Ray tracing pseudocode

```
// Compute camera settings  
// Clear image RGB buffer & instantiate sampler  
  
for (x = 0; x < width; ++x) {  
    for (y = 0; y < height; ++y) {  
        (px, py) = getPixelLoc(x, y, width, height);  
        dir = toWorld(px, py);  
        ray = Ray(camera.o, dir);  
        pixelColor = integrator->render(ray, sampler);  
        splatOnScreen(x, y, pixelColor);  
    }  
}
```

Offline rendering loop

Ray tracing pseudocode

```
// Compute camera settings  
// Clear image RGB buffer & instantiate sampler  
  
for (x = 0; x < width; ++x) {  
    for (y = 0; y < height; ++y) {  
        (px, py) = getPixelLoc(x, y, width, height);  
        dir = toWorld(px, py);  
        ray = Ray(camera.o, dir);  
        pixelColor = integrator->render(ray, sampler);  
        splatOnScreen(x, y, pixelColor);  
    }  
}
```

Generating rays

- ▶ Different ways to sample ray through pixel:
 - Single ray through pixel center (A1.1)
 - Single ray through random position on pixel
 - Multiple rays through pixel center
 - Multiple rays through random position on pixel (A1.2, *supersampling*)



Real-time rendering loop



Image source: Naughty Dog's Uncharted 4

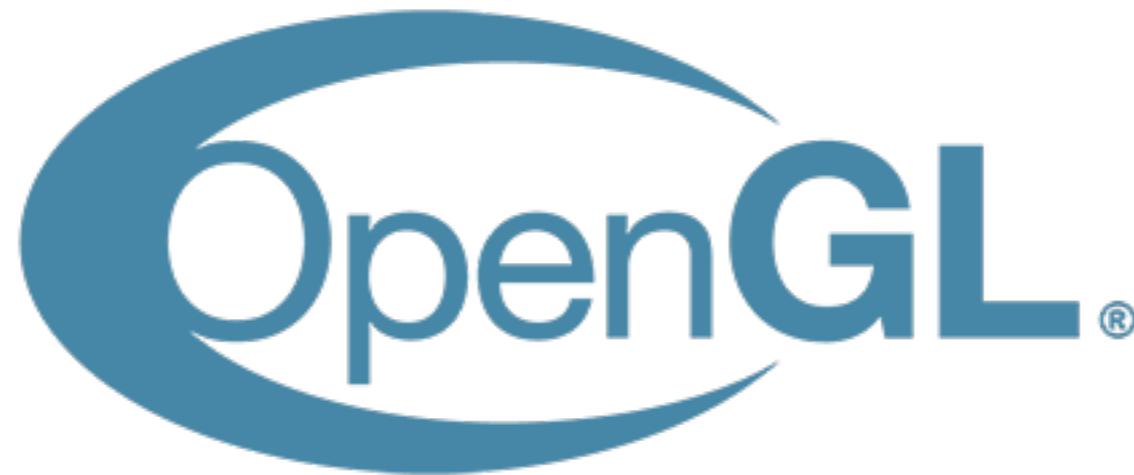
OpenGL 101

- ▶ OpenGL is an API to access the GPU
- ▶ This course uses a relatively modern version (3.3)
- ▶ Good references:
 - Official docs: <http://docs.gl>
 - Shaderific: <http://www.shaderific.com/glsl>

Real-time rendering loop

OpenGL 101

- ▶ GL = State Machine + Buffers
- ▶ Data is passed through *buffers*
- ▶ Commands are then used to change the *state* of the machine



Modern GL pipeline

1. Initialize buffers
2. For each frame:
 - a) Bind buffers (tell GPU which VAO/VBO to use)
 - b) Update uniforms
 - c) Draw call
3. Clean up memory

GL Buffers

- ▶ **Screen buffer**

Where all the pixels are written

- ▶ **Vertex Buffer Object (VBO)**

Where vertex attributes (position, normal, colour, etc.) are stored

- ▶ **Vertex Array Object (VAO)**

- ▶ **Other buffers**

Framebuffer Object (FBO)

Buffer Texture

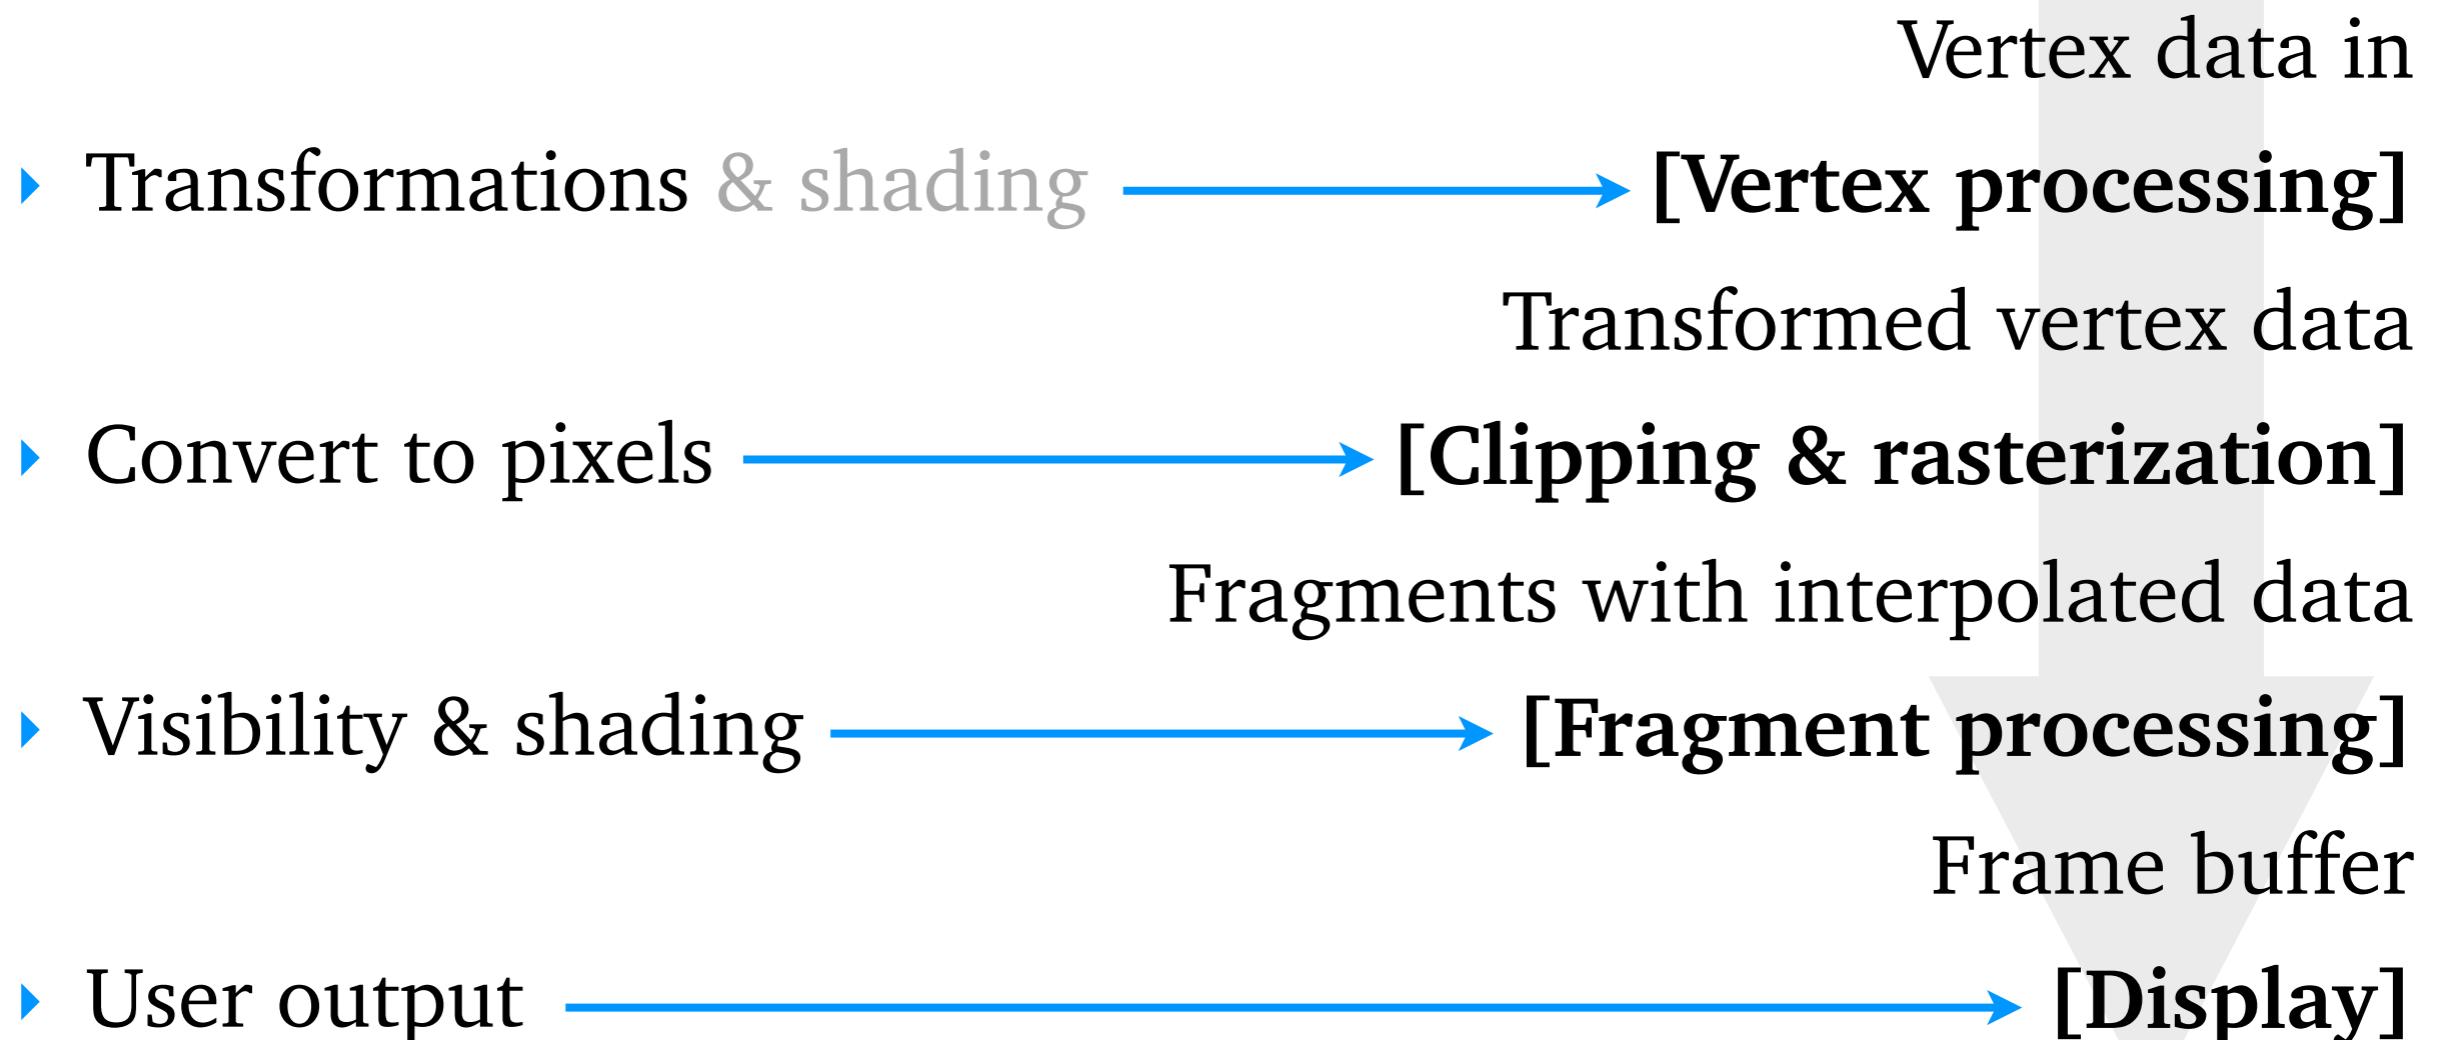
GL Buffers

- ▶ They are used the same way:
 1. Create buffer
 2. Bind buffer (to use it)
 3. Fill buffer with data

Rasterization (in one slide)

- ▶ Process of converting a continuous object to a discrete representation on a raster (pixel) grid
- ▶ At a high-level:
 1. Take triangles
 2. Figure out which pixels they cover
 3. Perform some computation on these pixels

Programmable rasterization pipeline



So, what do I code?

- ▶ Render pass that actually calls the shaders
- ▶ GLSL shaders:
 - Vertex shader (.vs file)
 - Fragment shader (.fs file)

Vertex shader

- ▶ Done at every vertex (hence the name)
- ▶ Means that it can manipulate attributes of vertices (position, normal, etc.)
- ▶ Comes earlier in the graphics pipeline

Fragment shader

- ▶ Comes *after* the vertex shader
- ▶ And *after* the rasterization step
- ▶ Takes care of how the pixels between the vertices look
- ▶ Interpolated between the defined vertices following specific rules to *fill in* the fragments

Shader logic

- ▶ Vertex shader will relay some information to the fragment shader (out)
- ▶ **Vertex shader**
Transform triangles to match camera view by setting `gl_Position`, and retrieve necessary information for shading
- ▶ **Fragment shader**
Compute the actual colour of the pixel using whatever info passed from the vertex shader

MVP Matrix

- ▶ Use GLSL uniforms:

```
gl_Position = projection * view * model  
             * vec4(position, 1.0);
```

- ▶ That's it. Field of view is handled by the code for the real-time part

TinyRender

General information

- ▶ Minimalistic rendering engine written in C++11
- ▶ Cross-platform (Windows, Linux and macOS)
- ▶ Both offline and real-time rendering frameworks
- ▶ Written by us (the TAs)
- ▶ Second year using it, hopefully bug-free™
- ▶ But, signal any important bug by email if you find one...

What you'll be able to render



Rendered with path tracer from A5

And with a few tweaks...

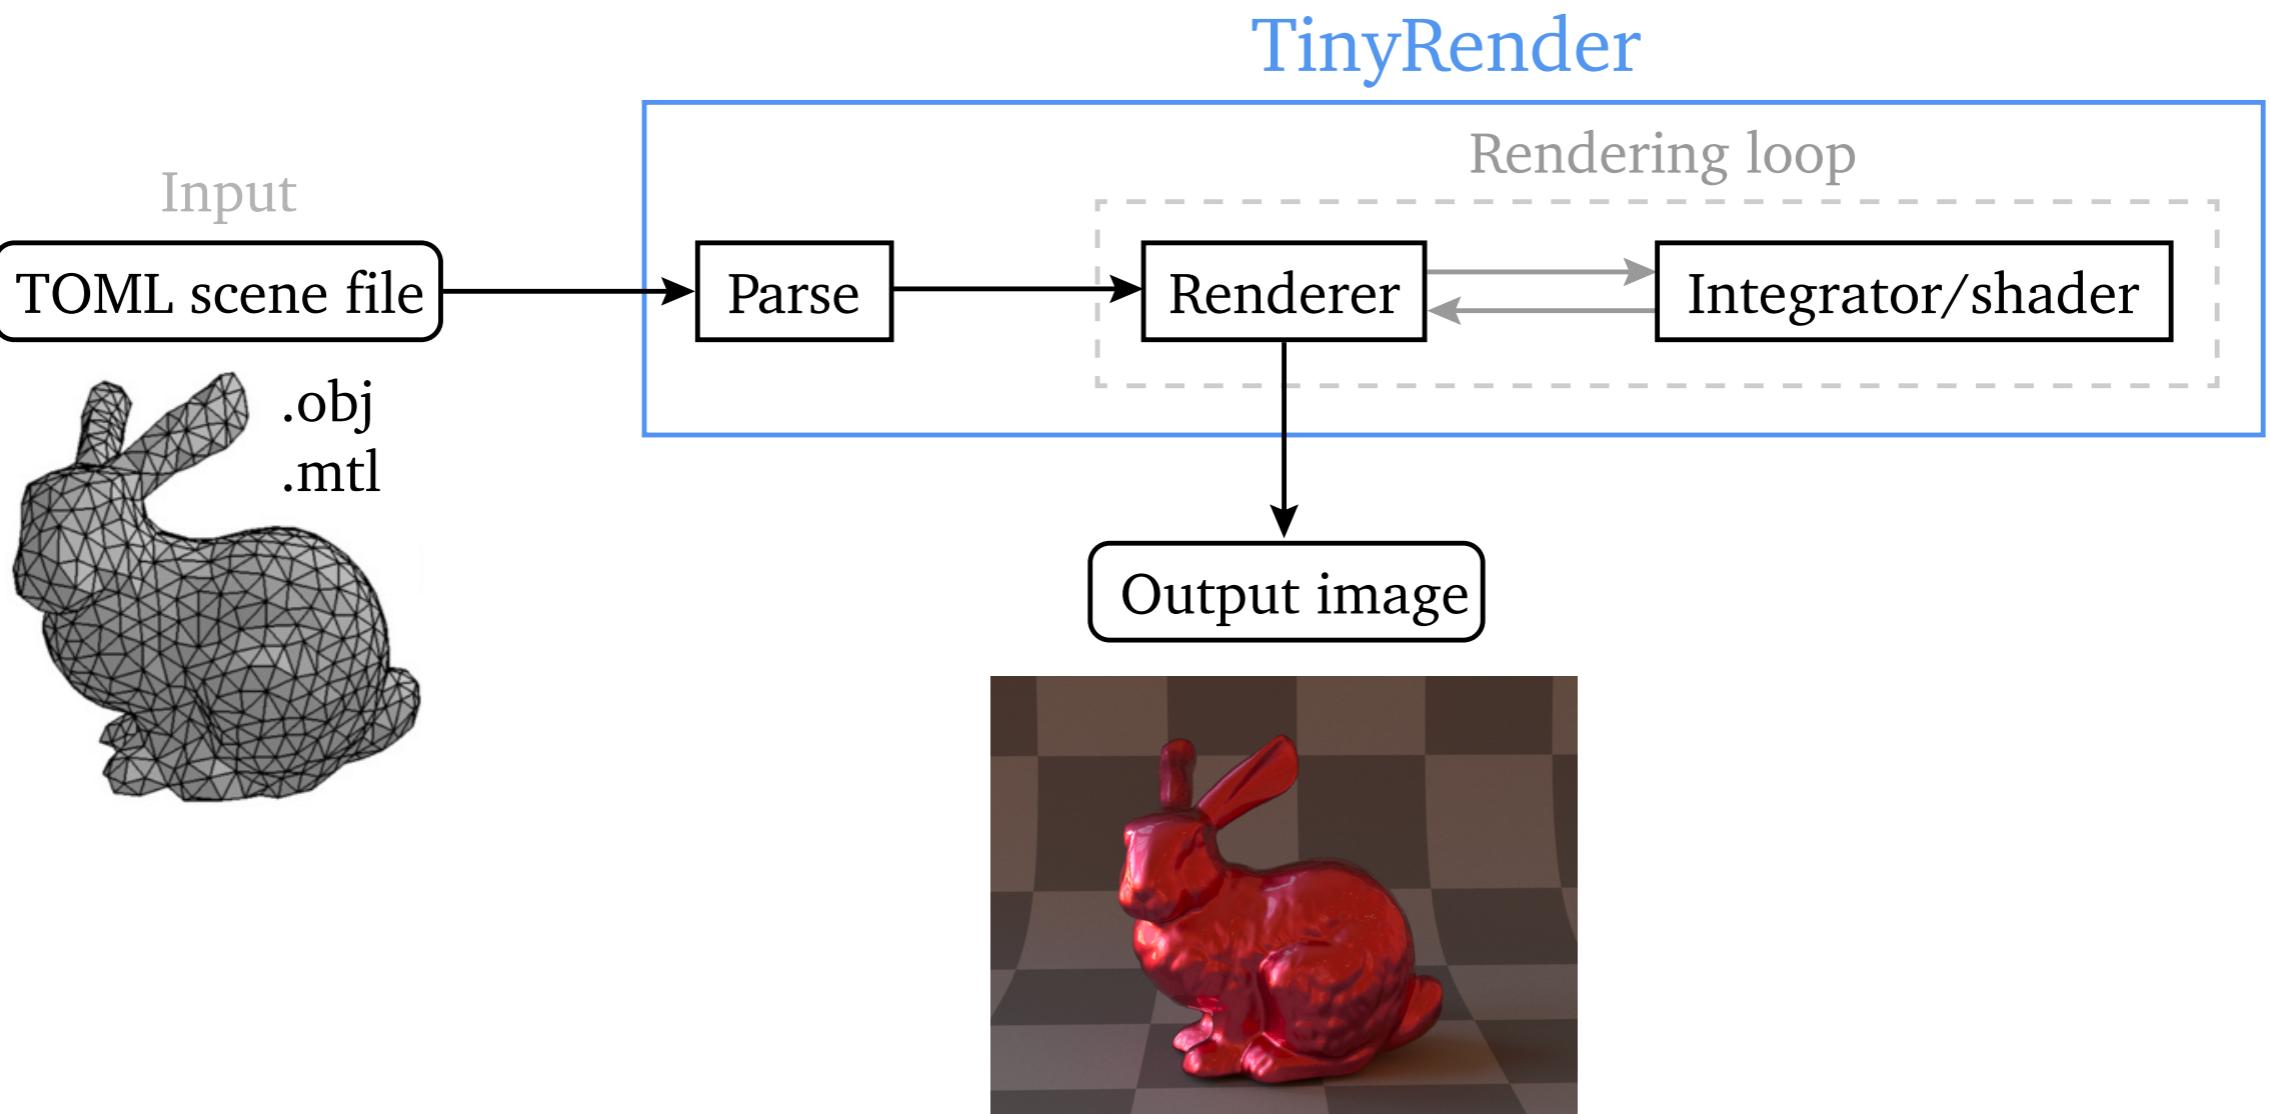


Gold



Brushed glass

Overview of the codebase



Codebase repository

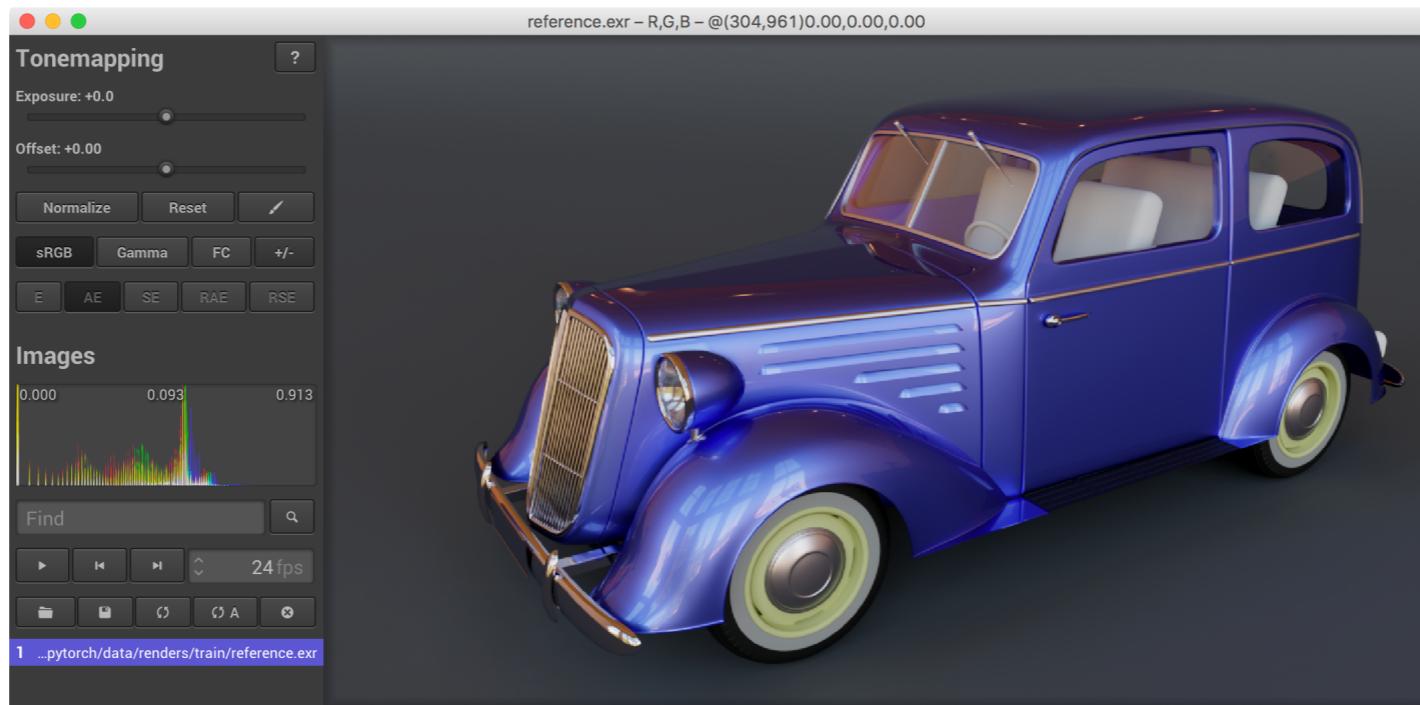
- ▶ Code available through a public Git repo
- ▶ Allows TAs to push modifications for upcoming assignments, bug fixes, etc. in a seamless way
- ▶ If you don't know Git, you will have to learn the very basics
- ▶ Only have to sync with master branch a handful of times throughout the semester

Building instructions

- ▶ All details can be found in the Read-me document
- ▶ Requires a few things to be installed beforehand
- ▶ If you want to compile TinyRender in the command-line under Ubuntu 14, you have every right to do so. **But**, we are *not* going to help you
- ▶ Do **NOT** upgrade to macOS Catalina when it comes out: may break the build and it will be **YOURS** to fix

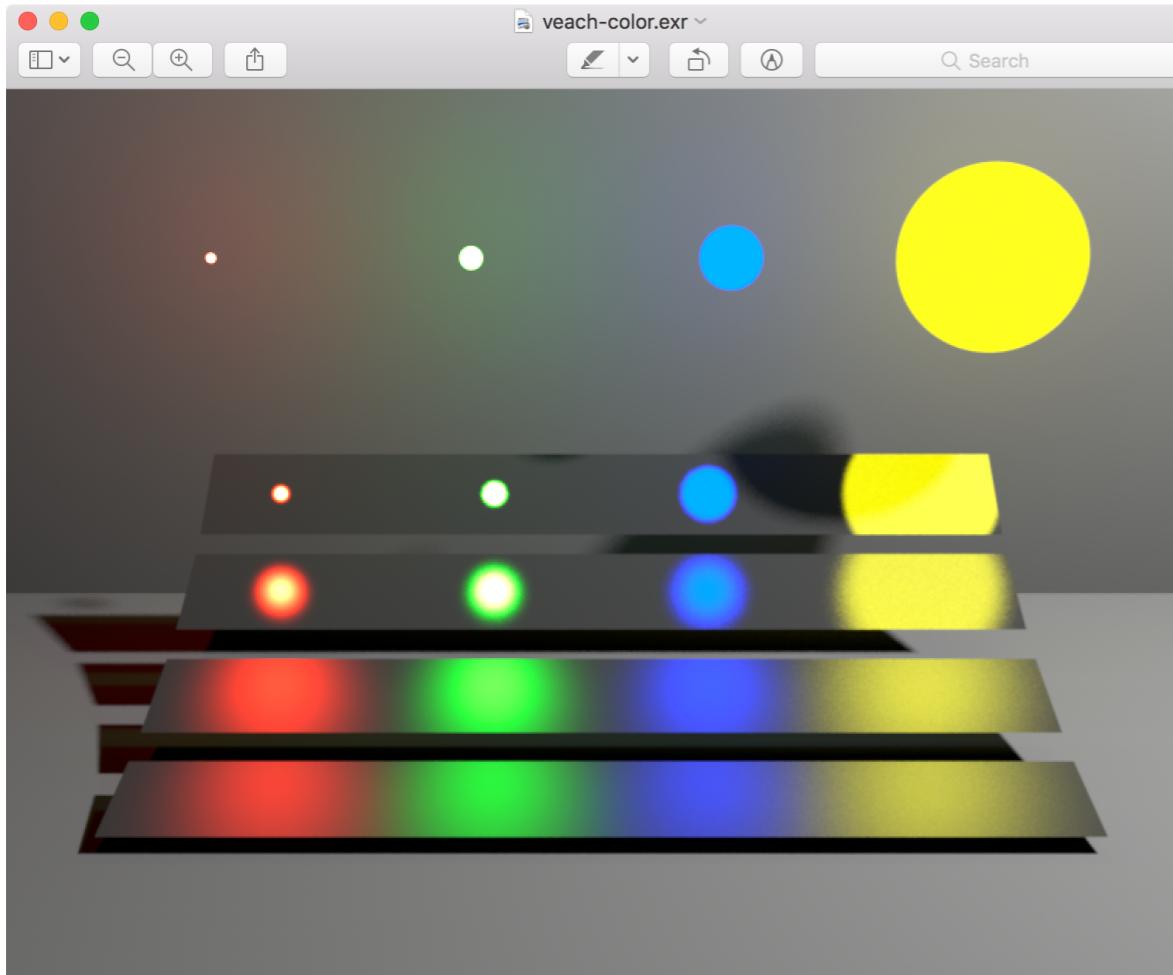
High dynamic range viewer

- ▶ Offline renders are saved in HDR formats (.exr)
- ▶ HDR = Pixels can take *float* values in $[0, +\infty)$, whereas standard LDR = Pixels are constrained to $[0,1]$
- ▶ Requires HDR viewer software (e.g. tev, HDRView)

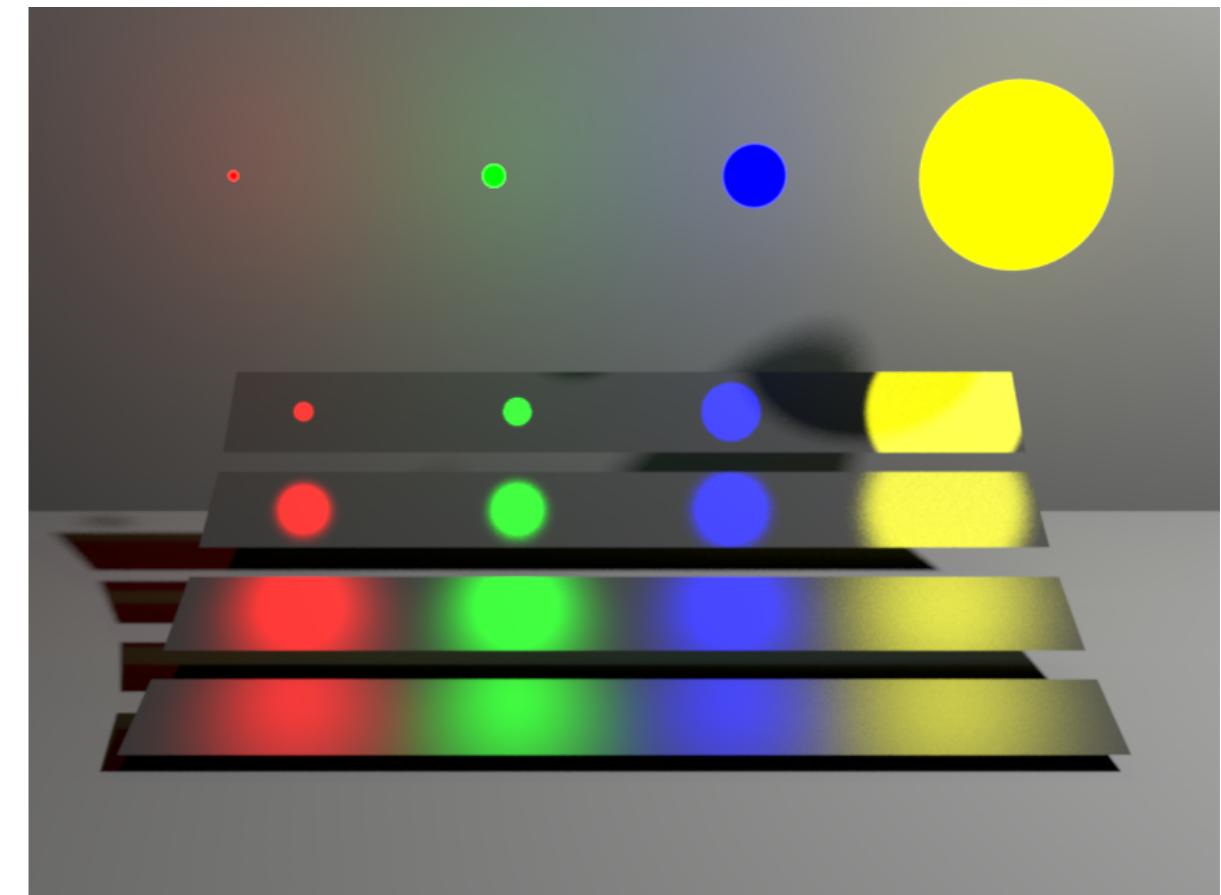


tev

Tonemapping and why it matters

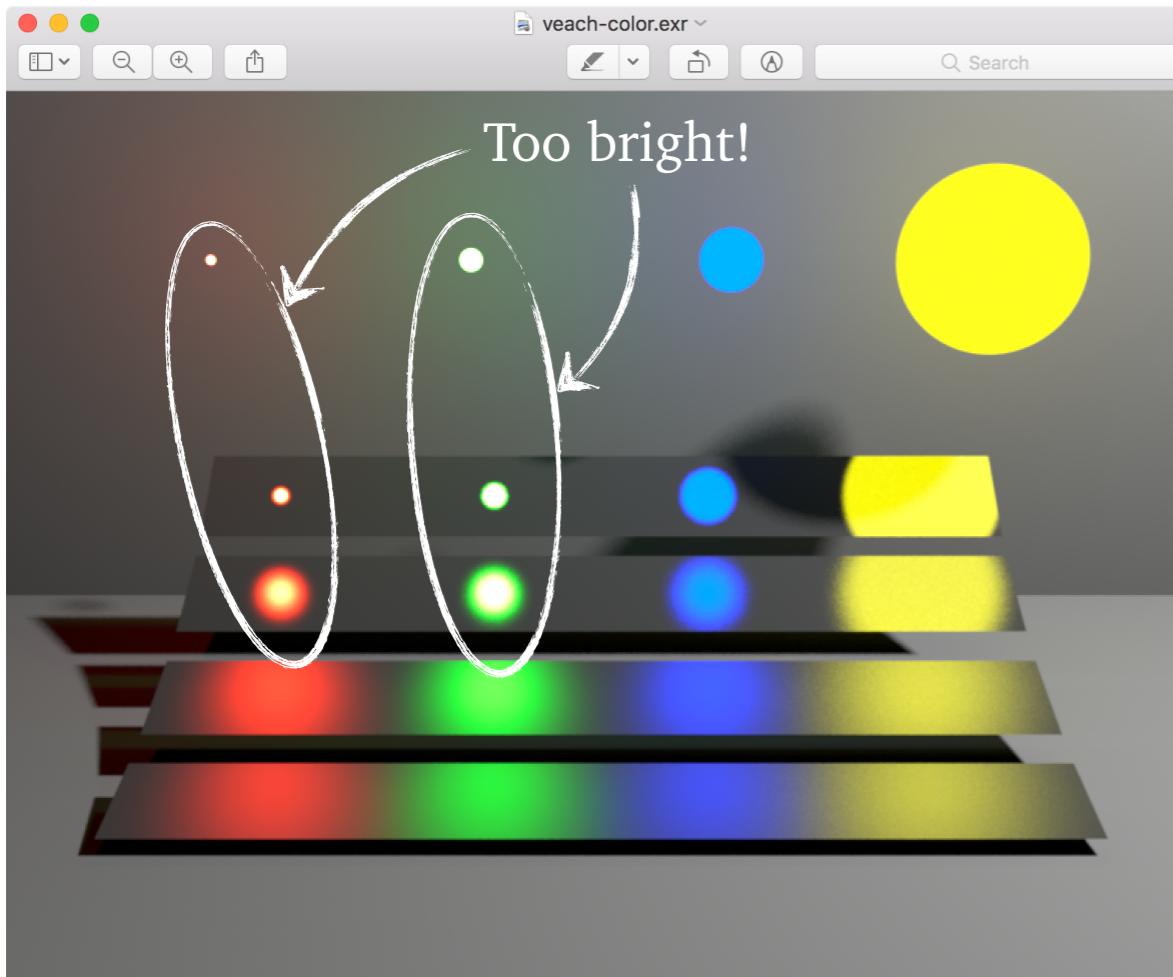


Wrong tonemap (Preview on macOS)

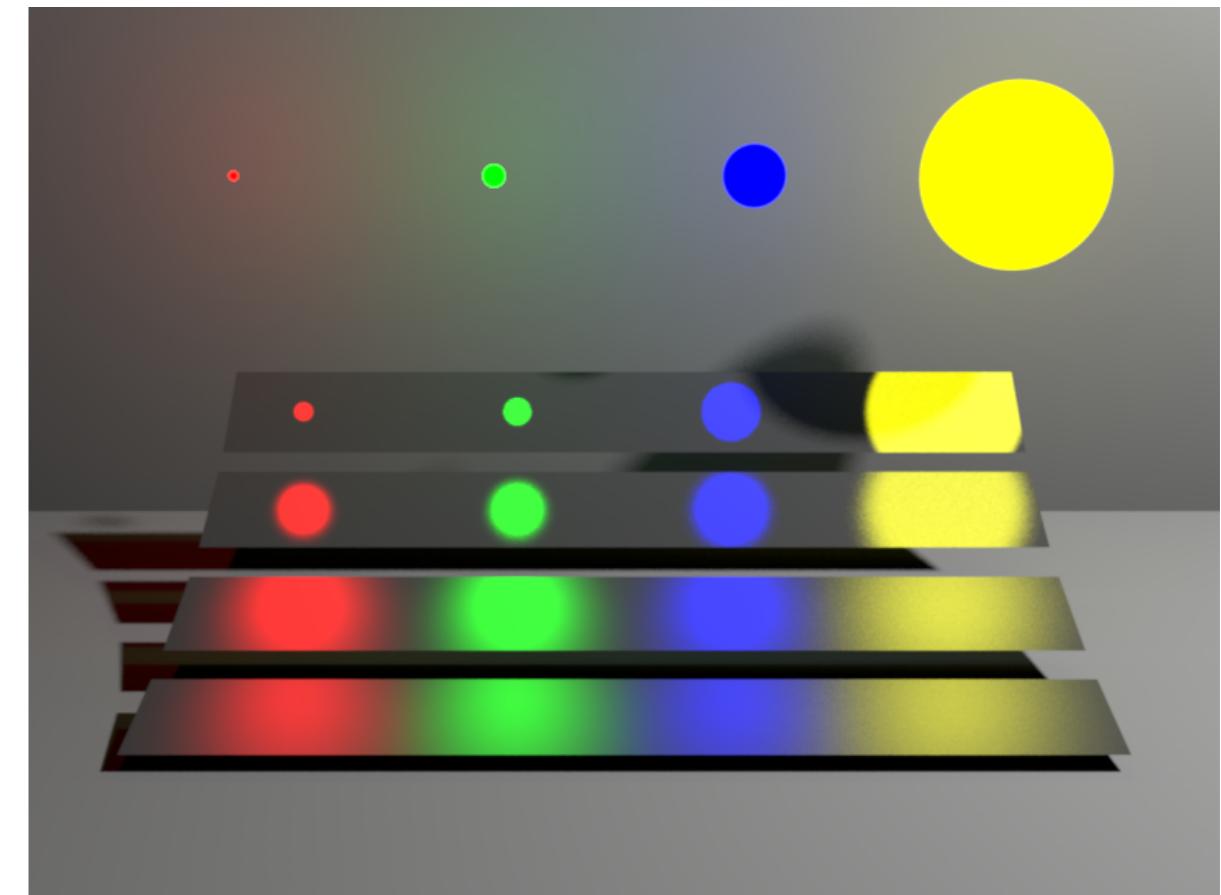


Correct tonemap

Tonemapping and why it matters



Wrong tonemap (Preview on macOS)



Correct tonemap

Important notes

- ▶ 3000+ lines of code: you will have to learn how to navigate the code base, *don't be scared*
- ▶ All assignments use TinyRender and you will have to recycle code across them
- ▶ Incremental level of difficulty, but some assignments are more important than others. **A1 is one of them**
- ▶ You are encouraged to discuss the assignments with your peers, but sharing code or images is *strictly forbidden*
- ▶ **This includes Github!**

Important notes

- ▶ Come to tutorials: TA's office hours = tutorials
- ▶ Use discussion board to ask questions
- ▶ Assignment deadlines **include** render times
- ▶ Submissions must follow a **strict** structure
- ▶ If you want to pursue grad studies in rendering, first step is doing well in this course!

Questions?